

## **STRESZCZENIE**

**Słowa kluczowe:**

**Dziedziny nauki i techniki zgodne z wymogami OECD:**

## **ABSTRACT**

**Keywords:**

**OECD field of science and technology:**

## CONTENTS

|   |    |
|---|----|
| List of abbreviations and nomenclature.....   | 7  |
| 1. Introduction .....   | 8  |
| 2. Theoretical background .....   | 10 |
| 2.1. Partially Observable Markov Decision Processes .....                                 | 10 |
| 2.2. Reinforcement Learning .....   | 11 |
| 2.3. Simulation-Based Search .....  | 16 |
| 2.4. Deep Learning .....  | 17 |
| 2.5. Variational Inference .....  | 21 |
| 3. State of the art .....   | 24 |
| 3.1. Learning and querying world models.....  | 24 |
| 3.2. Model-based reinforcement learning .....   | 26 |
| 3.2.1. World Models .....   | 26 |
| 3.2.2. Learning Latent Dynamics for Planning from Pixels .....                            | 29 |
| 3.2.3. Model-Based Reinforcement Learning for Atari .....                                 | 32 |
| 3.3. Monte-Carlo planning .....   | 33 |
| 3.3.1. A general reinforcement learning algorithm that masters chess, shogi, and Go ..... | 33 |
| 3.3.2. Value Prediction Network .....   | 34 |
| 4. Planning with learned model.....   | 36 |
| 4.1. HumbleRL framework .....   | 36 |
| 4.1.1. Architecture .....   | 36 |
| 4.2. Original World Models .....  | 38 |
| 4.2.1. World model .....  | 38 |
| 4.2.2. Controller.....  | 40 |
| 4.2.3. Data collection.....   | 40 |
| 4.2.4. Preprocessing .....  | 41 |
| 4.2.5. Implementation details .....   | 41 |
| 4.3. World Models and AlphaZero .....   | 42 |
| 4.3.1. World model .....  | 43 |
| 4.3.2. Controller.....  | 43 |
| 4.3.3. Implementation details, data collection and preprocessing.....                     | 44 |
| 4.4. PlaNet .....   | 45 |
| 4.4.1. World model .....  | 45 |
| 4.4.2. Planner.....   | 47 |
| 4.4.3. Data collection.....   | 48 |
| 4.4.4. Preprocessing .....  | 48 |
| 4.4.5. Implementation details .....   | 48 |
| 5. Experiments .....  | 49 |
| 5.1. Benchmarks.....  | 49 |

|   |    |
|---|----|
| 5.1.1. Arcade Learning Environment.....   | 49 |
| 5.1.2. Sokoban .....  | 51 |
| 5.2. Hardware .....   | 52 |
| 5.3. World Models for Sokoban .....   | 52 |
| 5.3.1. Train the unchanged World Models implementation in the Sokoban environment ..... | 53 |
| 5.3.2. Train World Models in Sokoban environment on 10x10 grid world states .....       | 55 |
| 5.3.3. Train World Model in Sokoban with auxiliary tasks .....                          | 56 |
| 5.4. World Models for Atari .....   | 57 |
| 5.4.1. Train unchanged World Models in the Boxing environment.....                      | 57 |
| 5.4.2. Train World Models with AlphaZero planner in the Boxing environment.....         | 59 |
| 5.5. PlaNet for Sokoban.....  | 60 |
| 5.5.1. Train unchanged PlaNet in the Sokoban environment.....                           | 60 |
| 5.6. PlaNet for Atari .....   | 61 |
| 5.6.1. Train unchanged PlaNet in the Boxing environment.....                            | 61 |
| 5.6.2. Train PlaNet in the Boxing environment with the increased action repeat .....    | 62 |
| 5.6.3. Train PlaNet in the Boxing environment with the lowered decoder variance ....    | 62 |
| 5.6.4. Train PlaNet in the Boxing environment with increased free nats .....            | 62 |
| 5.6.5. Train tuned PlaNet in the Freeway environment .....                              | 64 |
| 5.6.6. Train tuned PlaNet in the Freeway environment with a longer planning horizon     | 65 |
| 5.6.7. Train PlaNet in the Boxing environment without overshooting.....                 | 65 |
| 6. Conclusion .....   | 67 |
| References .....  | 68 |
| List of figures .....   | 71 |
| List of tables .....  | 72 |

## **LIST OF ABBREVIATIONS AND NOMENCLATURE**

## 1. INTRODUCTION

[\* *What is RL and model-free RL...*] Reinforcement learning, a subfield of artificial intelligence (AI), formalise rather the most obvious and common learning strategy among animals. It is learning how to achieve predefined goals through interaction with an environment [52]. Progress has been made in developing capable agents for numerous domains using deep neural networks in conjunction with model-free reinforcement learning [22][36][47], where raw observations directly map to agent's actions. However, current state-of-the-art approaches are very sample inefficient, they sometimes require tens or even hundreds of millions of interactions with the environment [35], and lack the behavioural flexibility of human intelligence, hence the resulting policies poorly generalize to novel tasks in the same environment.

[\* *Model-based RL with its benefits...*] The other branch of reinforcement learning algorithms, called model-based reinforcement learning, aims to address these shortcomings by endowing agents with a model of its environment. There are many ways of using the model: it can be used for data augmentation for model-free methods [12], some methods use the model to simulate experience for model-free methods to learn from [17], other methods focus on simulation-based search using the model [49] and there are even methods that integrate model-free and model-based approaches into the one architecture [55]. The model allows the agent to simulate an outcome of an action taken in a given state. The main upside is, that it allows the agent to plan by thinking ahead, seeing what would happen for a range of possible choices, and explicitly deciding between possible options without the risk of the adverse consequences of trial-and-error in the real environment - including making poor, irreversible decision. Even if the model needs to be synthesized from past real experience, it can be done more sample-efficiently than model-free methods, because it does not require propagating rewards through Bellman backups and can exploit additional unsupervised learning signals like future observations. Furthermore, the same model can be used by the agent to complete other tasks in the same environment [55] and planning carries the promise of increasing performance just by increasing the computational budget for searching for actions [50]. It gives AI hint of human intelligence flexibility and versatility.

[\* *Really briefly about what is problem of learning world dynamics of POMDP...*] Model-free methods are more popular and have been more extensively developed and tested than model-based methods. While model-free methods forego the potential gains in sample efficiency from using a model, they tend to be easier to implement and tune. The main downside of model-based reinforcement learning is that a ground-truth model of the environment is usually not available to an agent. If the agent wants to use a model in this case, it has to learn the model from experience, which creates several challenges. Because the model is often only an approximation of real environment, one of them is bias in the model that can be exploited by the agent [17], resulting in an agent which performs well with respect to the model, but behaves sub-optimally in the real environment. Different challenge also comes from fundamental downside of function approximation. The performance of agents employing common planning methods usually suffer from seemingly minor

model errors [53]. Those errors compound during planning, causing more and more inaccurate predictions the further horizon of a plan [54].

*[\* Long-term ambitious goal...]* There are many real-world problems that could benefit from application of general planning AI system. Company called DeepMind, driven by their experience from creating winning Go search algorithm AlphaZero [49], published AlphaFold [11], a system that predicts protein structure. The 3D models of proteins that AlphaFold generates are far more accurate than any that have come before making significant progress on one of the core challenges in biology.

Real-world applications of AI algorithms like this are often limited by the problem of sample inefficiency. In a setting with e.g. a physical robot the AI agent can not afford much trial-and-error behaviour, that could cause damage to the robot. It also can not afford running for hundreds of millions of time steps, for each task separately, in order to build a sufficiently large training dataset. Those machines work in the real world, not in an accelerated and parallelised computer simulation, and often need a human assistance. Progress in sample-efficient model-based algorithms is required in order to bring reinforcement learning into the real world applications.

*[\* Explain your topic...]* The aim of this work is to derive from previous work on model-learning in complex high-dimensional decision making problems [6][34][5][20] and apply them to planning in complex tasks. Those methods proved to train accurate models, at least in short horizon, and should open a path for application of planning algorithms like AlphaZero [49] to i.e. Atari 2600 games, a platform used for evaluation of general competency in artificial intelligence [35]. It should allow for improvement in data efficiency without loss in performance compared to model-free methods. This work focuses on three benchmarks: an arcade game with dense rewards Boxing, a challenging environment with sparse rewards Freeway and a complex puzzle game Sokoban.

## 2. THEORETICAL BACKGROUND

### 2.1. Partially Observable Markov Decision Processes

Markov Decision Processes (MDPs) are a classical formalization of sequential decision making, where actions influence not just immediate rewards, but also subsequent situations (states) and through those future rewards. Therefore, MDPs involve delayed rewards and the need to trade-off these with an immediate reward. Partially observable Markov Decision Processes (POMDPs), which describe a more general class of problems, have one major difference, the full state of an environment is unknown. The environment is perceived through observations that provide only partial information about the state. A good example are Atari games. Individual frames often does not provide full information about the game's state which is held in the game's RAM.

In this work following definition of MDP is used: it consists of a set of hidden states  $S$ , a set of observations  $O$  and a set of actions  $A$ . The dynamics of the MDP, from any state  $s \in S$  and for any action  $a \in A$ , are determined by transition function,  $P_{ss'}^a = p(S_{t+1} = s' | S_t = s, A_t = a)$ , specifying the distribution over the next state  $s' \in S$ . A reward function,  $R_{ss'}^a = p(R_{t+1} | S_t = s, A_t = a, S_{t+1} = s')$ , specifies the distribution over rewards for a given state transition. Finally, as mentioned earlier, POMDP is perceived through partial observations specified via probability distribution  $P_s = p(O_t | S_t = s)$ . A fixed initial state  $s_0$  is assumed. In episodic MDPs, which this work considers, an environment terminates with probability 1 in one of distinguished terminal states,  $s_T \in S$ , after finite number of transitions  $T$ . A return  $G_t = \sum_{k=t+1}^T \gamma^{k-t-1} r_k$  is the total reward accumulated in that episode from time  $t$  until reaching the terminal state at time  $T$ .  $0 \leq \gamma \leq 1$  is a discount factor that trade-offs short-term rewards with long-term rewards. A policy,  $\pi(s, a) = p(A_T = a | S_t = s)$ , maps a state  $s$  to a probability distribution over actions. A state value function,  $V_\pi(s) = \mathbb{E}_\pi[G_t | s_t = s]$ , is the expected return from state  $s$  when following policy  $\pi$  where the expectation is over the distributions of the environment and the policy. An action value function,  $Q_\pi(s, a) = \mathbb{E}_\pi[G_t | s_t = s, a_t = a]$ , is the expected return after selecting action  $a$  in state  $s$ , often called state-action pair, and then following policy  $\pi$  where, again, the expectation is over the distributions of the environment and the policy. Optimal state value and action value functions are unique value functions that maximise the value of every state or state-action pair,  $V^*(s) = \max_\pi V_\pi(s), \forall s \in S$  and  $Q^*(s, a) = \max_\pi Q_\pi(s, a), \forall s \in S, a \in A$ . The two are related to each other by this equality:  $V^*(s) = \max_a Q^*(s, a)$ . An optimal policy  $\pi^*(s, a)$  is a policy that maximises the optimal action value function for every state in the MDP,  $\pi^*(s, a) = \operatorname{argmax}_a Q^*(s, a)$ .

Reinforcement learning assume underlying MDP or POMDP, but the dynamics, the reward function and the observations distribution are hidden from it. Consequently, these can not be used directly for planning, but one could learn them through interaction with the environment. This concept is called planning and learning in literature [52] and it is fundamental for model-based reinforcement learning.

## 2.2. Reinforcement Learning

Reinforcement learning (RL) is learning what to do, how to map situations to actions, so as to maximise an expected return. [52] This mapping is called a policy  $\pi$ . RL consists of an agent that, in order to learn a good policy, acts in an environment. The environment provides a response to each agent's action  $a$  that is interpreted and fed back to the agent. The response consists of: reward  $r$  that is used as a reinforcing signal and state, or observation,  $s$  that is used to condition the agent's next decision. Fig. 21 explains it in the diagram. Each action-response-interpretation sequence is called a step, or a transition. Multiple steps in order form a trajectory. An episode is a trajectory that starts in an initial state  $s_0$  and finishes in a terminal state  $s_T$ . After the terminal state, the environment is reset in order to start the next episode from scratch. Very often, RL agents need dozens and dozens of episodes to gather enough experience to learn the near optimal policy. In many cases the policy is an approximation of some kind to the optimal policy and hence it will never be exactly optimal.

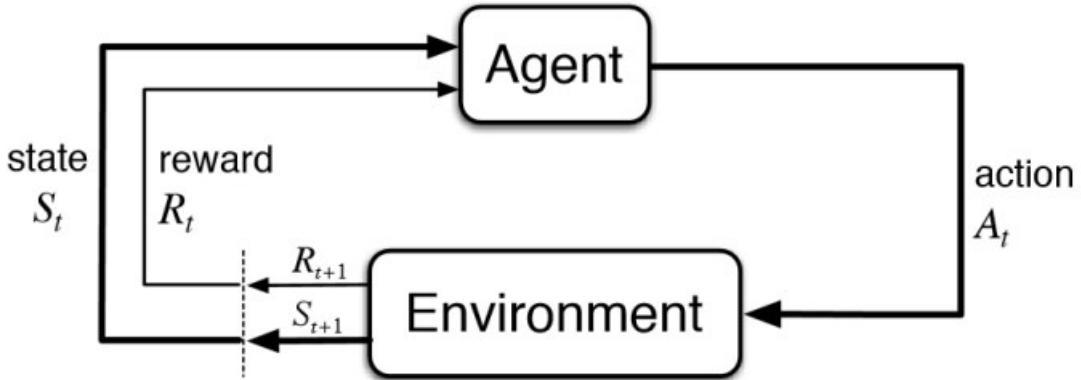


Fig. 21. Reinforcement Learning [52]

Each environment can be assigned one of two characteristics from 7 categories:

- Fully observable vs. partially observable: If an agent have access to the complete state of the environment at each point in time, then we say that the environment is fully observable. An environment might be partially observable because of noisy and inaccurate observations or because parts of the state are simply not accessible for the agent e.g. an automated taxi cannot see what other drivers are thinking.
- Single agent vs. multi-agent: The distinction between single-agent and multiagent environments may seem simple enough. For example, an agent solving a crossword puzzle by itself is clearly in a single-agent environment, whereas an agent playing chess is in a two-agent environment.
- Deterministic vs. stochastic: If the next state of the environment is completely determined by the current state and the action executed by the agent, then we say the environment is deterministic. Otherwise, it is stochastic. If the environment is partially observable, however, then it could appear to be stochastic. Most real situations are so complex that it is impossible

to keep track of all the unobserved aspects. For practical purposes, they must be treated as stochastic. It is often said that an environment is uncertain if it is not fully observable or not deterministic, because an agent can not be sure the outcomes of its actions.

- Episodic vs. continuing: In an episodic environment, the agent's experience is divided into episodes which finish at the terminal state. Crucially, the next episode does not depend on the actions taken in the previous episode. In continuing environments, on the other hand, the agent's experience may go on and on forever and there is no a distinct terminal state.
- Static vs. dynamic: If the environment can change while an agent is deliberating, then we say the environment is dynamic for that agent. Otherwise, it is static.
- Discrete vs. continuous: The discrete/continuous distinction applies to the internal state of the environment, to the way time is handled, and to the observations and actions of the agent. For example, the chess environment has a finite number of distinct states. Chess also has a discrete set of actions. Taxi driving is a continuous-state and continuous-time problem: the speed and location of the taxi and of the other vehicles sweep through a range of continuous values and do so smoothly over time. Taxi-driving actions are also continuous: steering angles, etc.
- Known vs. unknown: Strictly speaking, this distinction refers not to the environment itself but to the agent's (or designer's) state of knowledge about the "laws of physics" of the environment. In a known environment, the outcomes, or outcome probabilities if the environment is stochastic, for all actions are given. It means that the agent have access to the environment's MDP. Obviously, if the environment is unknown, the agent will have to learn how it works in order to make good decisions.

The term dynamic programming refers to a collection of algorithms that can be used to compute optimal policies given a perfect model of the environment as a Markov decision process (MDP). Classical dynamic programming algorithms are of limited utility in reinforcement learning both because of their assumption of a perfect model and because of their great computational expense, but they are still important theoretically. Dynamic programming provides an essential foundation for the understanding of the methods presented in this thesis. In fact, all of these reinforcement learning methods can be viewed as attempts to achieve much the same effect as dynamic programming, only with less computation and without assuming a perfect model of the environment.

First question to answer is: how to compute the state value function  $V(s)$  for an arbitrary policy  $\pi$ ? The process of doing so is called policy evaluation. It is easy to note that for all states  $s$ :  $V_\pi(s) = \mathbb{E}_\pi[G_t|s_t = s] = \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1}|s_t = s] = \mathbb{E}_\pi[R_{t+1} + \gamma V_\pi(s_{t+1})|s_t = s]$ , where the expectations are over the environment dynamics and are subscripted by  $\pi$  to indicate that they are conditional on  $\pi$  being followed. The existence and uniqueness of  $V_\pi$  are guaranteed as long as either  $\gamma < 1$  or eventual termination is guaranteed from all states under the policy  $\pi$ . If the environment's dynamics are completely known, then the equation is a system of  $|S|$ , number of states, simultaneous linear equations in  $|S|$  unknowns. In principle, its solution is a straightforward,

if tedious, computation. For more practical purposes, iterative solution methods are available [52], but not described here.

Major reason for computing the value function for a policy is to help find better policies. With determined the value function  $V_\pi$  for an arbitrary deterministic policy  $\pi$ , the question now is: if for some state  $s$  changing the policy to deterministically choose an action  $a \neq \pi(s)$ , different then from the current policy, yields higher expected return? The value function tells how good it is to follow the current policy from  $s$ , but would it be better or worse to change to the new policy? One way to answer this question is to consider selecting  $a$  in  $s$  and thereafter following the existing policy  $\pi$ . The value of this way of behaving is  $Q_\pi(s, a) = E[R_{t+1} + \gamma V_\pi(s_{t+1})|s_t = s, a_t = a]$ , where expectation is over the environment dynamics. The key criterion is whether this is greater than or less than  $V_\pi(s)$ . If it is greater — that is, if it is better to select  $a$  once in  $s$  and thereafter follow  $\pi$  than it would be to follow  $\pi$  all the time — then one would expect it to be better still to select  $a$  every time  $s$  is encountered, and that the new policy would in fact be a better one overall. It is a natural extension to consider changes at all states and to all possible actions, selecting at each state the action that appears best according to  $Q_\pi(s, a)$  and this way create a new, improved policy:  $\pi' = \underset{a}{\operatorname{argmax}} Q_\pi(s, a)$ . The process of making a new policy that improves on an original policy, by making it greedy with respect to the value function of the original policy, is called policy improvement. If the new policy is as good as, but not better then, the old policy, then it is the optimal policy, which can be proved without much trouble [52].

Policy iteration consists of two simultaneous, interacting processes, one making the value function consistent with the current policy, policy evaluation, and the other making the policy greedy with respect to the current value function, policy improvement. In policy iteration, these two processes alternate, each completing before the other begins, but this is not really necessary. In value iteration, for example, only a single iteration of policy evaluation is performed in between each policy improvement. In asynchronous dynamic programming methods, the evaluation and improvement processes are interleaved at an even finer grain. In some cases a single state is updated in one process before returning to the other. As long as both processes continue to update all states, the ultimate result is typically the same: convergence to the optimal value function and an optimal policy.

The term generalized policy iteration is used to refer to the general idea of letting policy-evaluation and policy-improvement processes interact, independent of the granularity and other details of the two processes. Almost all reinforcement learning methods are well described as general policy iteration. That is, all have identifiable policies and value functions, with the policy always being improved with respect to the value function and the value function always being driven toward the value function for the improvement policy, as suggested by the fig.22. If both the evaluation process and the improvement process stabilize, that is, no longer produce changes, then the value function and policy must be optimal [52].

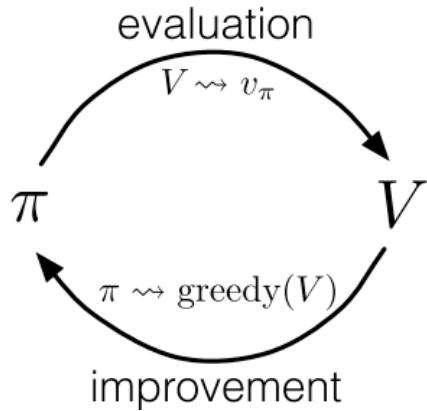


Fig. 22. General Policy Iteration [52]

Reinforcement learning faces two dilemmas. The first one is: How do you distribute credit for success among the many decisions that may have been involved in producing it? It is called credit-assignment problem, which past actions to reinforce for a positive outcome.

The second one is exploration-exploitation trade-off. To find a good policy, such that maximise an expected return, an agent needs to explore its options: if it would behave greedily all the time, it simply could not know if other actions lead to better returns. On the other hand, the agent also needs to exploit its knowledge about an environment in order to do well and progress in the environment to otherwise unavailable states. In other words, for policy evaluation to work, continual exploration needs to be assured. The most common approach to assuring that all state-action pairs are encountered is to consider only policies that are stochastic with a nonzero probability of selecting all actions in each state.

So called exploration policies are made to exploit an agent's knowledge about an environment, but at the same time, these constantly explore the environment's state-space. These policies are told to be soft, meaning that  $\pi(a|s) > 0$  for every action in every state. One such family of policies is called  $\epsilon$ -greedy, meaning that most of the time they choose an action that has maximal estimated action value, but with probability  $\epsilon$  they instead select an action at random. That is, all non-greedy actions are given the minimal probability of selection  $\frac{\epsilon}{|A(s)|}$ , where  $|A(s)|$  is number of legal actions in state  $s$ , and the remaining bulk of the probability,  $1 - \epsilon + \frac{\epsilon}{|A(s)|}$ , is given to the greedy action. The bigger the  $\epsilon$ , the higher the exploration rate and vice versa.

There are two general approaches to reinforcement learning: on-policy methods and off-policy methods. On-policy methods attempt to evaluate or improve the policy that is used to make decisions, whereas off-policy methods evaluate or improve a policy different from that used to generate the data. In the former case, an agent is extra sensitive to imperfect updates, in cases when its value function or its policy are somehow approximated. Because the agent learns directly from its behaviour, if the bad update makes the agent behave worse than before, the agent effectively falls back in its learning progress. Also, the on-policy methods learn action values not for the optimal policy, but for a near-optimal policy that still explores, which is required as noted before.

An example of on-policy algorithm would be Monte-Carlo control [52]. Because it is reinforcement learning setting an agent do not have access to underlying MDP or POMDP and needs to learn from interactions with an environment, from its experience. An obvious way to an action value function learning, or policy evaluation, is simply to play with the environment using the current policy and average the returns observed after visits to each state-action pair. It is called Monte-Carlo prediction. As more returns are observed, the average should converge to the expected value. Policy improvement is done by making the policy  $\epsilon$ -greedy with respect to the current value function. The two proceed according to the idea of generalized policy iteration.

In off-policy case, the behavioural exploratory policy is used for experience collection, which is then used to steadily improve the target policy towards the optimum. These methods, although favourable, are often more complex. Because the data is due to a different policy, off-policy methods are often of greater variance and are slower to converge. On the other hand, off-policy methods are more powerful and general. They include on-policy methods as the special case in which the target and behavior policies are the same. Off-policy methods also have a variety of additional uses in applications. For example, they can often be applied to learn from data generated by a conventional non-learning controller, or from a human expert. An example of off-policy algorithm would be Q-Learning [52], which is not described here.

There are two major kinds of reinforcement learning methods, model-free and model-based. Model-based methods rely on planning as their primary component, while model-free methods primarily rely on learning. The word planning is used in several different ways in different fields. In reinforcement learning the term is used to refer to any computational process that takes a model as input and produces or improves a policy for interacting with the modeled environment. Although there are real differences between these two kinds of methods, there are also great similarities. In model-free reinforcement learning, the agent samples episodes of real experience and updates its value function from real experience. In model-based reinforcement learning the agent samples episodes of simulated experience and updates its value function from simulated experience. This symmetry between learning and planning has an important consequence: algorithms for learning can also become algorithms for planning, simply by substituting simulated experience in place of real experience.

The model of an environment can mean anything that an agent can use to predict how the environment will respond to its actions. Given a state and an action, a model produces a prediction of the resultant next state and next reward. If the model is stochastic, then there are several possible next states and next rewards, each with some probability of occurring. Some models produce a description of all possibilities and their probabilities. These are called distribution models. Other models produce just one of the possibilities, sampled according to the probabilities. These are called sample models. The kind of model assumed in dynamic programming, estimates of the MDP's dynamics  $P_{ss'}^a$  and  $R_{ss'}^a$ , is a distribution model.

Models can be used to mimic or simulate experience. Given a starting state and action, a sample model produces a possible transition, and a distribution model generates all possible transitions

weighted by their probabilities of occurring. Given a starting state and a policy, a sample model could produce an entire episode, and a distribution model could generate all possible episodes and their probabilities. In either case, it is said that the model is used to simulate the environment or rollout the trajectory and produce simulated experience.

In artificial intelligence, there are two distinct approaches to planning according to the definition presented. State-space planning, which is viewed primarily as a search through the state-space for an optimal policy or an optimal path to a goal. Actions cause transitions from state to state, and value functions are computed over states.

In what is called plan-space planning, planning is instead a search through the space of plans. Operators transform one plan into another, and value functions, if any, are defined over the space of plans. Plan-space planning includes evolutionary methods.

Within a planning agent, there are at least two roles for real experience: it can be used to improve the model, to make it more accurately match the real environment, and it can be used to directly improve the value function and policy using the kinds of model-free reinforcement learning. The former is called model-learning, and the latter is called direct reinforcement learning. The possible relationships between experience, model, values, and policy are summarized in the fig.23. It can be seen, that experience can improve value functions and policies either directly or indirectly via the model.

Planning is often conducted in small, incremental steps. This enables planning to be interrupted at any time, which appears to be a key requirement for efficiently intermixing planning with acting and with learning of the model. Planning in very small steps may be the most efficient approach even in pure planning problems if the problem is too large to be solved exactly or the agent can not wait for exact solution and needs to act based on approximated one.

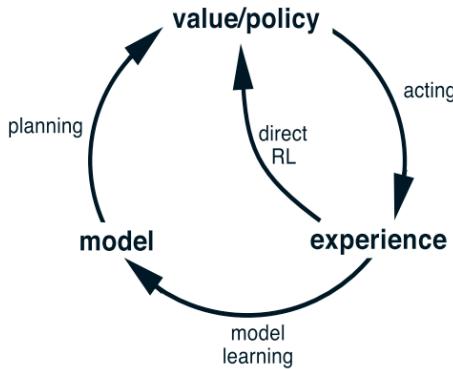


Fig. 23. Model-based Reinforcement Learning. [52] Each arrow shows a relationship of influence and presumed improvement.

### 2.3. Simulation-Based Search

Simulation-based search is something called planning at decision time [52] and can be used, of course, in model-based reinforcement learning setting. It requires a sample model of the MDP, which can sample state transitions and rewards from  $P_{ss'}^a$  and  $R_{ss'}^a$ , respectively. However,

it is not necessary to know these probability distributions explicitly. The next state and reward could be generated by a black box simulator. The efficiency and effectiveness of simulation-based search depends in large part on the performance and accuracy of the model. The model learning problem is described later. Now, access to a perfect sample model is assumed.

Simulation-based search algorithms sample experience in sequential episodes. Each simulation begins in a root state  $s_0$ . At each step  $u$  of simulation, an action  $a_u$  is selected according to a simulation policy, and a new state  $s_{u+1}$  and reward  $r_{u+1}$  is generated by the sample model. This process repeats, without backtracking, until a terminal state is reached. The values of states or actions are then updated from the simulated experience.

Simulation-based search is usually applied online, at every time-step  $t$ , by initiating a new search that starts from the current state  $s_0 = s_t$ . The distribution of simulations then represents a probability distribution over future experience from time-step  $t$  onwards. Simulation-based search exploits temporality, in other words focuses on the current situation, by learning from this specific distribution of future experience, rather than learning from the distribution of all possible experience. Furthermore, as the agent's policy improves, the future experience distribution will become more refined. It can focus its value function on what is likely to happen next, given the improved policy, rather than learning about all possible eventualities.

Monte-Carlo simulation is a very simple simulation-based search algorithm for evaluating candidate actions from a root state  $s_0$ . The search proceeds by simulating complete episodes from  $s_0$  until termination, using a fixed simulation policy. The action values  $Q(s_0, a)$  are estimated by the mean outcome of all simulations with candidate action  $a$ .

Monte-Carlo tree search (MCTS) is perhaps the best-known example of a simulation-based search algorithm. It makes use of Monte-Carlo simulation to evaluate the nodes of a search tree. There is one node,  $n(s)$ , corresponding to every state  $s$  encountered during simulation. Each node contains a total count of visits of the state,  $N(s)$ , a value  $Q(s, a)$  and a visit count  $N(s, a)$  for every possible action in a state  $s$ . Simulations start from the root state  $s_0$ , and are divided into two stages. When state  $s_u$  is contained in the search tree, a tree policy selects the action with the highest value  $Q(s, a)$ . Otherwise, a random default policy is used to roll out simulations to completion. After each simulation,  $s_0, a_0, r_1, s_1, a_1, \dots, r_T$ , each node  $n(s_u)$  in the search tree is updated incrementally to maintain the count and mean return from that node,

$$N(s_u) \leftarrow N(s_u) + 1$$

$$N(s_u, a_u) \leftarrow N(s_u, a_u) + 1$$

$$Q(s_u, a_u) \leftarrow Q(s_u, a_u) + \frac{G_u - Q(s_u, a_u)}{N(s_u, a_u)}$$

## 2.4. Deep Learning

Machine learning gives AI systems the ability to acquire their own knowledge, by extracting patterns from raw data. It stands in opposition to classical computer programs which execute

explicit instructions hand-coded by a programmer. One example of machine learning algorithm is logistic regression. It can determine whether to recommend cesarean delivery or not [38]. Another widely used machine learning algorithm called naive Bayes can distinguish between legitimate and spam e-mail.

The performance of these machine learning algorithms depends heavily on the representation of the problem they are given. For example, when logistic regression is used to recommend cesarean delivery, the AI system does not examine the patient's MRI scan directly. It would not be able to make useful predictions as individual pixels in an MRI scan have negligible correlation with any complications that might occur during delivery. It, instead, gets several pieces of relevant information, such as the presence or absence of a uterine scar, from the doctor. Each piece of information included in the representation of the data is known as a feature. Logistic regression learns the relation between those features and various outcomes, such as a recommendation of cesarean delivery. The algorithm does not influence the way that the features are defined in any way.

Sometimes it can be hard to hand-craft a good problem's representation. For example, suppose that someone would like to write a program to detect cats in photographs. People know that cats are furry and have whiskers, so they might like to use the presence of a fur and whiskers as features. Unfortunately, it is difficult to describe exactly what a fur or a whisker looks like in terms of pixel values. This gets even more complicated when taking into account e.g. shadows falling on the cat or an object in the foreground obscuring part of the animal. One solution to this problem is to use machine learning to discover not only the mapping from representation to output but also the representation itself. This approach is known as representation learning. Learned representations often result in much better performance of machine learning algorithms than can be obtained with hand-designed representations. They also allow AI systems to rapidly adapt to new tasks with minimal human intervention.

Deep learning is a particular kind of machine learning that achieves great power and flexibility by learning to represent the world as a nested hierarchy of concepts, with each concept defined in relation to simpler concepts. Deep learning solves representation learning problem by introducing representations that are expressed in terms of other, simpler representations. Fig. 24 shows how a deep learning system can represent the concept of an image of a person by combining simpler concepts, such as corners and contours, which are in turn defined in terms of edges.

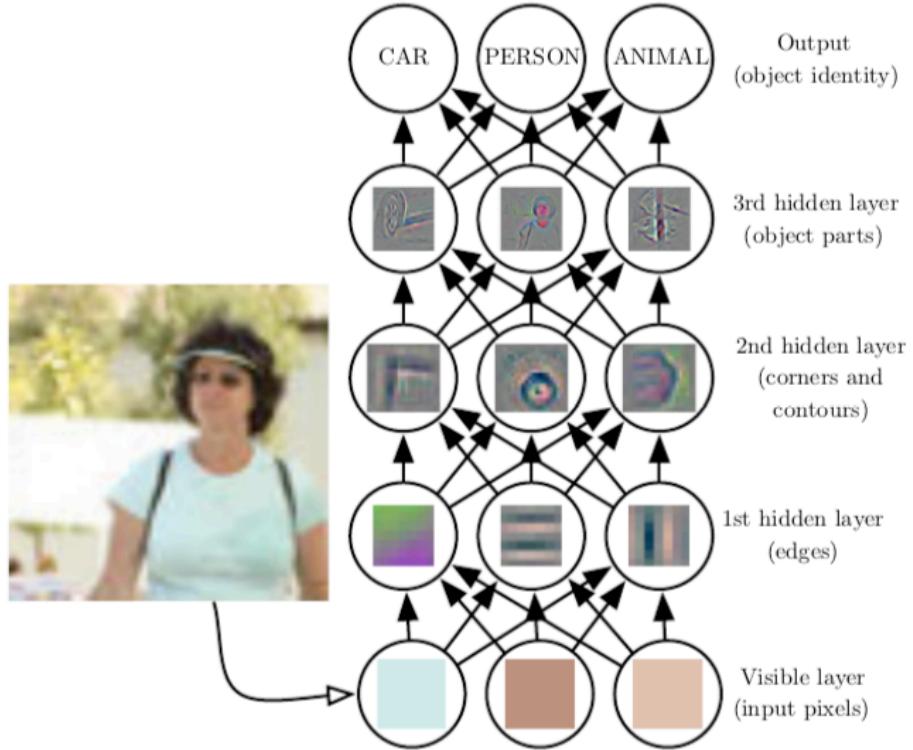


Fig. 24. Deep Learning [13]

The fundamental example of a deep learning model is a fully-connected (FC) neural network (NN), sometimes called a multilayer perceptron (MLP). A MLP is just a mathematical function mapping some set of input values to output values. The function is formed by composing many simpler functions, called neurons or perceptrons, into layers. Each layer provides a new representation of its input, which might be a vector, to all the neurons in the next layer by multiplying the input element-wise with its parameters and applying some kind of non-linear activation function to the sum of these products e.g. sigmoid function. Without the mentioned activation function, the multilayer NN would just simplify to a linear mapping without ability to model non-linear relationships of inputs and outputs of the NN. Fig. 25 shows example MLP and dependencies between perceptrons. The input is presented at the input layer. Then a hidden layer (or series of them) extracts increasingly abstract features from the image. These layers are called “hidden” because their values are not given in the data. Instead, the model must determine which concepts are useful for explaining the relationships in the observed data. Finally, this description of the input in terms of the features can be used to produce the output at the output layer.

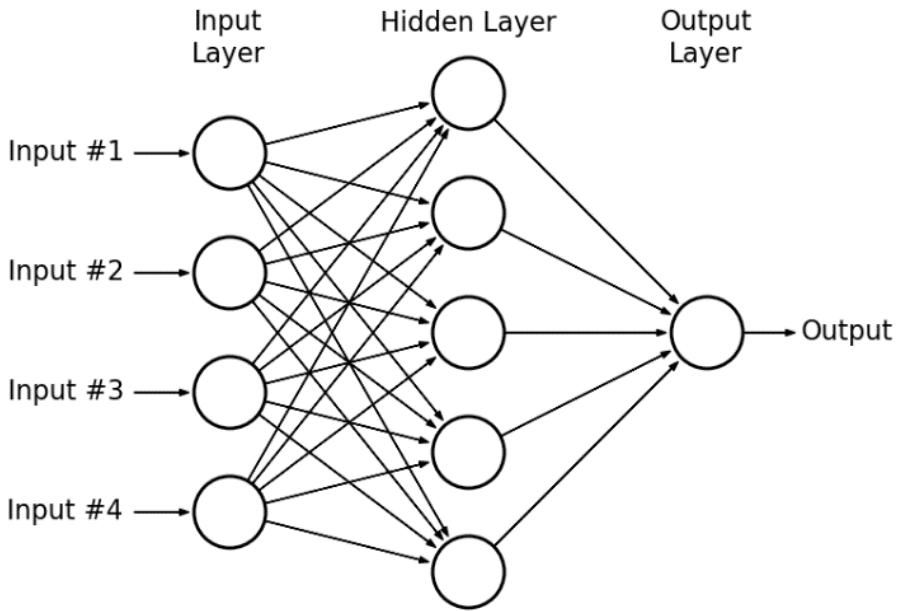


Fig. 25. Multilayer perceptron

Other type of very common neural network is a recurrent neural network (RNN). Humans do not start their thinking from scratch every second. As they read this thesis for example, they understand each word based on their understanding of previous words. Their thoughts have persistence. MLPs can not do this, that is why RNNs address this issue. These are networks with loops in them, allowing information to persist. In the fig. 26, a chunk of neural network,  $A$ , looks at some input  $x_t$  and outputs a value  $h_t$ . A loop allows information to be passed from one step of the network to the next. It turns out that RNNs are not all that different than a normal feedforward neural network. A recurrent neural network can be thought of as multiple copies of the same network, each passing a message to a successor.

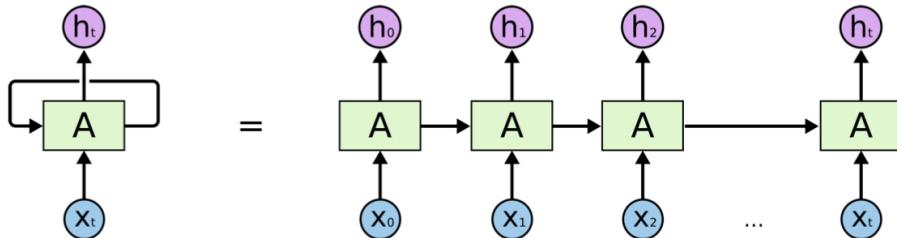


Fig. 26. A recurrent neural network [44]

Yet another very useful NN architecture is a convolutional neural network (CNN) which makes use of the fact that inputs to a NN can often have similar patterns in different parts of the input. For example in images, the object can appear in different positions in the image, but it is still the same object. CNNs exploit that fact by grouping parameters into filters in each layer and slide, or more precisely convolve, each filter across the width and height of the input volume and compute dot products between the entries of the filter and the input at any position. The 2D outputs

of all filters are stacked together to form a 3D activation map. It is then processed by the next, similar, layer with its own set of filters. At the end of CNN, the final 3D volume is often flattened into a vector and then further processed by a fully-connected NN. The example CNN is shown in fig. 27.

A CNN is able to successfully capture the spatial dependencies in an image through the application of relevant filters. The architecture performs a better fitting to the data with spatial hierarchy due to the reduction in the number of parameters involved and reusability of weights.

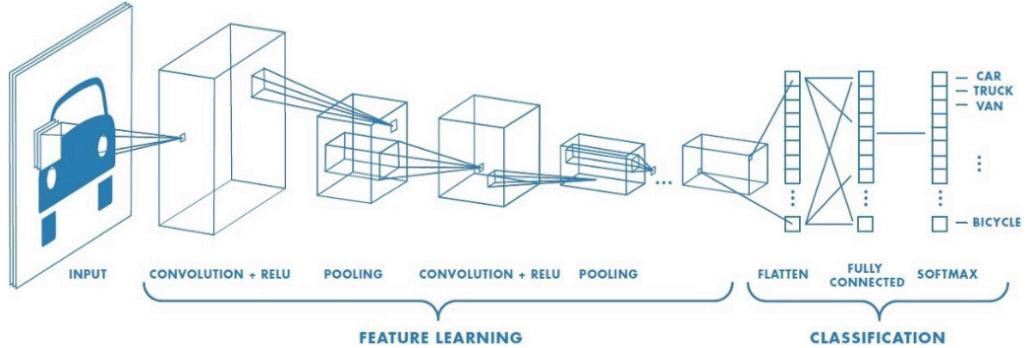


Fig. 27. A convolutional neural network [43]. ReLU [40] and Softmax are activation functions. Pooling means that each filter's output is divided into, for instance,  $2 \times 2$  regions and only maximum value from each region is passed to the next layer.

A neural network is trained, in the simplest case, to output a correct target value for each input from a dataset. This is done by backpropagating some measure of error, how much different is the NN's output value from the target value, through the NN. This produces gradients. Gradients are vectors that point in the direction of steepest ascent of the NN's error. These are used in the opposite direction to iteratively change parameters and lower the error on the dataset. The process is repeated iteratively, updates are small steps, because gradients are first order derivatives and the directions these point in are accurate only locally. This algorithm is called gradient decent (GD) and it is at the heart of the deep learning. More about neural networks and training procedures can be found in the Deep Learning book [13].

## 2.5. Variational Inference

Generative modeling is a broad area of machine learning which deals with models of distributions  $p(X)$ , defined over datapoints  $X$  in some potentially high-dimensional space. For instance, images are a popular kind of data for which one might create generative models. Each datapoint, which is an image, has thousands or millions of dimensions, in form of pixels, and the generative model's job is to somehow capture the dependencies between pixels. For example, that nearby pixels have similar color, and are organized into objects. Then, one often cares about producing more examples, that are like those already in a database, using the trained generative model. When training a generative model, the more complicated the dependencies between the dimensions, the more difficult the models are to train. For instance, the problem of generating images of handwritten digits from 0 to 9. If the left half of the character contains the left half of a 5, then the

right half cannot contain the left half of a 0, or the character will very clearly not look like any real digit. Intuitively, it helps if the model first decides which character to generate before it assigns a value to any specific pixel. This kind of decision is formally called a latent variable or a code. That is, before the model draws anything, it first randomly samples a digit value  $z$  from the set  $[0, \dots, 9]$ , and then makes sure all the strokes match that character. It is called latent, because it is not observed, or included in the dataset, and must be inferred in order to discover it.

To make this notion precise mathematically, the aim is to maximise the probability of each datapoint  $X$  in the training set under the entire generative process according to:  $p(X) = \int p(X|z)p(z)dz$ , where  $p(X|z)$  is a likelihood distribution, which generates images from latent variables, and  $p(z)$  is a prior distribution, which decides on latent variables. Calculating  $p(X)$  for most interesting, non-linear, models is computationally intractable. One way to bypass this is to use variational inference.

From now on, POMDP notation will be used: datapoints are now observations  $o$  and latent variables are now states  $s$ . Variational inference can be used to train latent variable models by maximizing a lower bound of the marginal probability density  $p(o) = \int p(o|s)p(s)ds$ . The term  $p(o|s)$  is the likelihood of the data given the latent variables  $s$  and, in the case of neural networks, takes the form of a parameterised model. Instead of dealing with the prior  $p(s)$  directly, variational autoencoders (VAEs) infer  $p(s)$  using the posterior  $p(s|o)$  [28]. As the true form of the posterior distribution  $p(s|o)$  is unknown, variational inference turns the problem of inferring latent variables into an optimisation problem by approximating the true posterior distribution with a variational distribution  $q(s|o)$ , called an encoder, which takes the form of a simpler distribution such as a fully factorized Gaussian, meaning the latent state vector dimensions are independent from each other, parameterised by a neural network and then minimising the Kullback-Leibler (KL) divergence between  $q(s|o)$  and  $p(s)$ . As the KL divergence is nonnegative and minimised when  $q$  is the same as  $p$ , the training objective for VAEs is known as the variational or evidence lower bound (ELBO) on the data log-likelihood:

$$\ln p(o) \geq \mathbb{E}_{q(s|o)} [\ln p(o|s)] - D_{KL}[q(s|o)||p(s)]$$

where  $p(s)$  is very often a standard normal distribution to simplify calculations and because no prior knowledge about latent variables is assumed. The left term of ELBO is called reconstruction loss, as it makes an observation  $o$  more probable given its latent variable  $s$ . The right term of ELBO is called regularisation loss or complexity or information bottleneck, as it pulls the approximate posterior  $q$  closer to the uninformed prior  $p$  and ensures this way that it explains the observation in the simplest possible way, preventing overfitting and aiding generalisation. VAEs parameters are trained using stochastic gradient descent, as common in the deep learning setting, to maximise the variational bound and this way indirectly maximise the probability of the dataset under the generative model.

Because the POMDP experience forms trajectories, something called structured variational inference is used. The joint probability of states  $s$ , observations  $o$ , and rewards  $r$  conditioned on the actions  $a$  from the initial timestep  $t = 1$  to the end of a trajectory of length  $T$  is:

$$p(o_{1:T}, r_{1:T}, s_{1:T} | a_{1:T}) = p(o_1 | s_1) p(s_1) \prod_{t=2}^T p(o_t | s_t) p(r_t | s_{t-1}, a_{t-1}) p(s_t | s_{t-1}, a_{t-1})$$

where, besides the prior for the initial state  $p(s_t)$ , a transition model prior  $p(s_t | s_{t-1}, a_{t-1})$  is introduced.

Since the model is non-linear, it is not possible to directly compute the state posteriors that are needed for parameter learning. Instead, an encoder  $q(s_{1:T} | o_{1:T}, a_{1:T}) = \prod_{t=1}^T q(s_t | s_{t-1}, a_{t-1}, o_t)$  is used to infer approximate state posteriors from past observations and actions, where  $q(s_t | s_{t-1}, a_{t-1}, o_t)$  could be a diagonal Gaussian with mean and variance parameterised by a neural network. Here, the filtering posterior [33] is used that conditions on past observations since, at the end, the model is designed to simulate future states and rewards based on past observations and actions. One may also use the full smoothing posterior during training [15], where the posterior for a single latent variable  $s_t$  depends on all future observations, rewards and actions too, but it could not be used in RL setting where at each timestep the future is unknown.

Using the encoder, a variational bound on the data log-likelihood can be constructed. For simplicity, here only losses for predicting the observations are written, but the reward losses follow by analogy:

$$\ln p(o_{1:T} | a_{1:T}) \geq \sum_{t=1}^T \left( \mathbb{E}_{q(s_t | o_{\leq t}, a_{<t})} [\ln p(o_t | s_t)] - \mathbb{E}_{q(s_{t-1} | o_{\leq t-1}, a_{<t-1})} [D_{KL}[q(s_t | o_{\leq t}, a_{<t}) || p(s_t | s_{t-1}, a_{t-1})]] \right)$$

Here, the left term of ELBO is still a reconstruction loss, but now the right term of ELBO trains the approximate posterior and the transition model to be consistent with each other.

### 3. STATE OF THE ART

Quite surprisingly, there is not much work on model-based planning and learning in complex environments, like video games, from images. This chapter researches related work that helps arrive at the final solution of this thesis' problem of planning in imagination.

#### 3.1. Learning and querying world models

In this paper [5], the authors aim to address the challenge of learning accurate, computationally efficient models of complex domains, a key challenge in model-based reinforcement learning, and using them to solve RL problems. First, they advocate the use of computationally efficient state-space environment models that make predictions at a higher level of abstraction, both spatially and temporally, than at the level of raw pixel observations. Such models substantially reduce the amount of computation required to make predictions, as future states can be represented much more compactly. Second, in order to increase model accuracy, they examine the benefits of explicitly modeling uncertainty in state transitions.

In the fig. 31, there are different environment models described in the paper. The models,  $p$ , are learned in an unsupervised way from observations,  $o$ , conditioned on actions,  $a$ . In particular, the authors focus on how fast and accurately models can predict, at time step  $t$ , some future statistics  $x_{t+1:t+\tau}$ , e.g. an environment's rewards, over a horizon  $\tau$  by doing Monte-Carlo rollouts of the model given an arbitrary sequence of actions  $a_{t:t+\tau-1}$ , that can later be used for decision making.

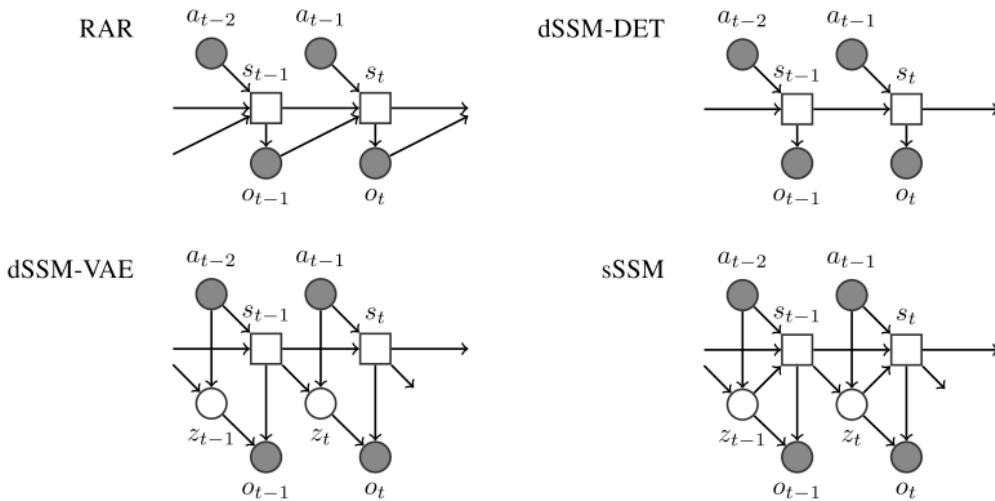


Fig. 31. The graphical models representing the architectures of different environment models. Boxes are deterministic nodes, circles are random variables and filled circles represent variables observed during training [5].

A straight-forward choice is the family of temporally auto-regressive models over the observations,  $o_{t+1:t+\tau}$ . If these use a recurrent mapping that recursively updates sufficient statistics  $s_t = f(s_{t-1}, a_{t-1}, o_{t-1})$ , therefore reusing the previously computed statistics  $s_{t-1}$ , and use them to predict future observations,  $p(o_{t+1:t+\tau} | o_{\leq t}, a_{<t+\tau}) = \prod_{r=t+1}^{t+\tau} p(o_r | f(s_{r-1}, a_{r-1}, o_{r-1}))$ , then the au-

thors call these models recurrent auto-regressive models (RAR). If  $f$  is parameterised as a neural network, RARs are equivalent to recurrent neural networks. Although faster then auto-regressive models that don't reuse statistics from previous time steps, these are still expect to be slow when generating Monte-Carlo rollouts, as they still need to explicitly render observations  $o_{t+1:t+\tau}$  in order to make any predictions,  $x_{t+1:t+\tau}$ .

State-space models (SSMs) circumvent this by positing that there is a compact state representation  $s_t$  that captures all essential aspects of the environment on an abstract level. It is assumed that  $s_{t+1}$  can be predicted from the previous state  $s_t$  and action  $a_t$  alone, without the help of previous pixels  $o_{\leq t}$ ,  $p(s_{t+1}|s_{\leq t}, a_{\leq t}, o_{\leq t}) = p(s_{t+1}|s_t, a_t)$ . Furthermore, it is assumed that  $s_t$  is sufficient to predict  $o_t$ , i.e.  $p(o_t|s_{\leq t}, a_{\leq t}) = p(o_t|s_t)$ . This modelling choice implies that the latent states are, by construction, sufficient to generate any future predictions  $x_{t+1:t+\tau}$ . Hence, the model never have to directly sample costly high-dimensional pixel observations.

The authors consider two flavors of SSMs: deterministic SSMs (dSSMs) and stochastic SSMs (sSSMs). For dSSMs, the latent transition  $s_{t+1} = f(s_t, a_t)$  is a deterministic function of the past, whereas for sSSMs, they consider transition distributions  $p(s_{t+1}|s_t, a_t)$  that explicitly model uncertainty over the next state. The sSSMs are parameterised by introducing for every  $t$  a latent variable  $z_t$  whose distribution depends on  $s_{t-1}$  and  $a_{t-1}$ , and by making the state a deterministic function of the past state, action, and latent variable:  $z_{t+1} \sim p(z_{t+1}|s_t, a_t)$ ,  $s_{t+1} = f(s_t, a_t, z_{t+1})$ . The observation model computes the conditional distribution  $p(o_t|\cdot)$ . It either takes as input the state  $s_t$  (deterministic decoder), or the state  $s_t$  and latent  $z_t$  (stochastic decoder). sSSMs always use the stochastic decoder. dSSMs can use either the deterministic decoder (dSSM-DET), or the stochastic decoder (dSSM-VAE). The latter can capture joint uncertainty over pixels in a given observation  $o_t$ , but not across time steps. The former is a fully deterministic model, incapable of modeling joint uncertainties (in time or in space).

The paper provides the first comparison of deterministic and stochastic, pixel-space and state-space models w.r.t. speed and accuracy, applied to challenging environments from the Arcade Learning Environment [2]. Specifically, the prediction of dSSM-DET exhibits "sprite splitting", or layering of multiple possible future observations, at corridors of MS-PACMAN [56], whereas multiple samples from the sSSM show that the model has a reasonable and consistent representation of uncertainty in this situation. Moreover, the authors note that SSMs, which avoid computing pixel renderings at each rollout step, exhibit a speedup of more then 5 times over the standard AR model.

Extensive experiments establish that state-space models accurately capture the dynamics of Atari games from raw pixels. The computational speed-up of state-space models, while maintaining high accuracy, makes their application in RL feasible. The authors demonstrate that agents which query these models for decision making outperform strong model-free baselines on the game MS-PACMAN, demonstrating the potential of using learned environment models for planning.

The authors explore the topic of learning fast generative models that can be used in model-based reinforcement learning for planning. Their ideas find application in PlaNet [20] and World Models [17], the stochastic state-space model (sSSM), and hence in this thesis' solution too.

### **3.2. Model-based reinforcement learning**

Model-free reinforcement learning, although can be used to learn effective policies for complex tasks, typically requires very large amounts of data. In fact, substantially more than a human would need to learn the same games. How can people learn so quickly? Part of the answer may be that people can learn how the game works and predict which actions will lead to desirable outcomes. Similar mechanism is used by model-based reinforcement learning. Papers below show a current state of this approach to learning.

#### **3.2.1. World Models**

In World Models [17] paper, its authors explore the idea of using large and highly expressive neural networks, that can learn rich spatial and temporal representation of data, and applying them to reinforcement learning. The RL algorithm is often bottlenecked by the credit assignment problem, which makes it hard for traditional RL algorithms to learn millions of weights of a large model. To accomplish their goal, they decompose the problem of an agent training into two stages: they first train a generative neural network to learn a model of the agent's world in an unsupervised manner. Thereafter, by using a compressed spatial and temporal representation of the environment extracted from the world model as inputs to the agent, they train a linear model to learn to perform a task in the environment. The small linear model lets the training algorithm focus on the credit assignment problem on a small search space, while not sacrificing capacity and expressiveness via the larger world model.

Their solution consists of three components: Vision for encoding the spatial information, Memory for encoding the temporal information and Controller which represents the agent's policy. Fig. 32 depicts a flow diagram of the agent's model.

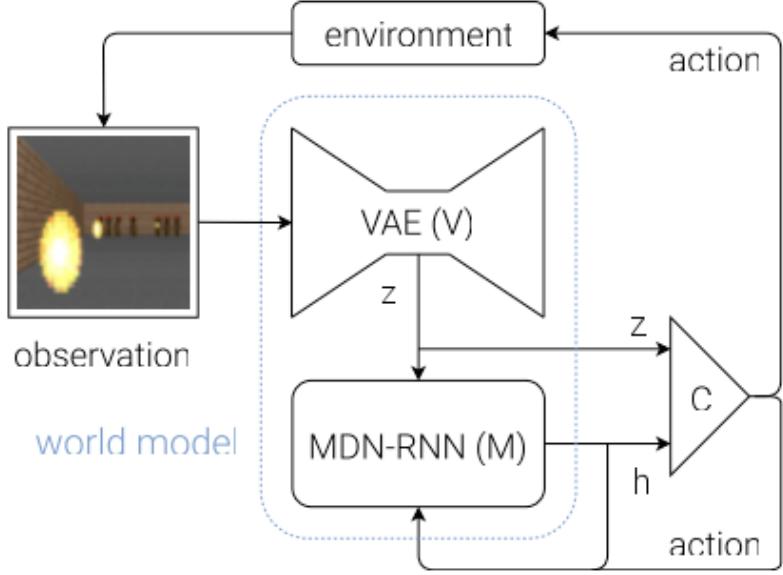


Fig. 32. Flow diagram of the agent’s model [17]. The raw observation is first processed by the Vision at each time step  $t$  to produce  $z_t$ . The input into the Controller is this latent vector  $z_t$  concatenated with the Memory hidden state  $h_t$  at each time step. The Controller will then output an action vector  $a_t$  and will affect the environment. The Memory will then take the current  $z_t$  and action  $a_t$  as an input to update its own hidden state to produce  $h_{t+1}$  to be used at time step  $t+1$ .

The environment provides the agent with a high dimensional visual observation, a game frame, at each time step. The essential task of the Vision model is to encode this high dimensional observation into a low dimensional latent state. To do this, Vision is implemented as Variational Autoencoder [28]. It is trained in an unsupervised manner on randomly generated experience from the environment. The authors assume that the random agent can efficiently explore environment and no iterative training procedure is implemented. The dataset is gathered once and fixed for Vision and Memory training.

Since many complex environments are partially observable, the visual observation at each time step, and hence the latent state, doesn’t include full information about the current situation in the environment. To acquire full knowledge, the agent needs to encode what happens over time. This is the role of the Memory. It is implemented as popular recurrent neural network (RNN) architecture called Long Short-Term Memory [24] and trained on the same data as Vision to predict the next step future latent state that Vision is expected to produce. Because many environments are stochastic in nature, the RNN is trained to output a probability density of the next latent state approximated as a mixture of Gaussian distribution - in literature, this approach is known as Mixture Density Network combined with a RNN [14] (MDN-RNN). Moreover, using the stochastic Memory the authors are able to train more robust Controller, more on that later.

To be more precise, the MDN-RNN will model  $p(z_{t+1}|o_{\leq t}, a_{\leq t}) = p(z_{t+1}|h_t) \prod_{i=1}^t q(z_i|o_i)$ , where  $h_t = f(h_{t-1}, z_t, a_t)$  is the hidden state of the RNN,  $f$ , that encodes past information about states and actions from the beginning of the episode until the time step  $t$ . Furthermore,  $o_t$ ,  $z_t$  and  $a_t$  are the observation, the latent state and the action at time step  $t$  respectively. During sampling the

authors adjust a temperature parameter  $\tau$ , that scale mixing coefficients in the MDN, to control model uncertainty [16]. They find it useful for training the Controller later on.

The Controller model represents the agent’s policy. It is responsible for determining course of actions to take in order to solve a given task. Controller is a simple linear model that maps the concatenated latent state  $z_t$  and hidden state  $h_t$  at the time step  $t$  directly to the action  $a_t$  at that time step:  $a_t = W[z_t h_t] + b$ , where  $W$  and  $b$  are the weight matrix and bias vector of that model. The authors deliberately made Controller as simple as possible, and trained it separately from Vision and Memory, so that most of the agent’s complexity resides in the world model (Vision and Memory). The latter can take the advantage of current advances in deep learning that provide tools to train large models efficiently when well-behaved and differentiable loss function can be defined. Shift in the agent’s complexity towards the world model allows the Controller model to stay small and focus its training on tackling the credit assignment problem in challenging RL tasks. It is trained using evolution strategy, which is rather an unconventional choice that only currently have been considered as a viable alternative to popular RL techniques [45].

Their solution is able to solve an OpenAI Gym’s CarRacing environment [4], which is the continuous-control, top-down racing task. It is the first known solution to achieve the score required to solve this task. Nonetheless, because the Controller uses real experience for training counted in millions of episodes, they did not improve sample-efficiency compared to other model-free solutions [26]. But, what is really interesting, in the process of training the Memory learns to simulate the original environment. The authors show that the learned Controller can function inside of the imagined environment of CarRacing, that is, simulated by the Memory. In the second experiment, they show that the agent can not only play in the imagination, but also it is able to learn solely from imagined experience, produced by its Memory, and successfully transfer this policy back to the actual environment of VizDoom (see fig. 33). In the first trials, the Controller learned to exploit imperfect simulations of the Model, which only approximates the true environment dynamics. To mitigate this behaviour, they adjust the temperature parameter of MDN-RNN to control the amount of randomness in the Memory, hence controlling the trade-off between realism and exploitability. The Controller learns from simulated experience, which means that only tens of thousands of episodes from the environment are needed for the training of Vision and Memory. Assuming that each episode consists of hundreds of frames, millions of frames in total are required for training. It makes World Models very sample efficient compared to other state-of-the-art model-free methods that require even two orders of magnitude more data [36].



Fig. 33. VizDoom: the agent must learn to avoid fireballs shot by monsters from the other side of the room with the sole intent of killing the agent [17].

The authors results indicate that their world model is able to model complex environments from visual observations and it can be used for planning. Therefore, it may prove useful for the topic of this thesis.

### 3.2.2. Learning Latent Dynamics for Planning from Pixels

The authors propose the Deep Planning Network [20] (PlaNet), a purely model-based agent that learns the environment dynamics from images and chooses actions through fast online planning in latent space. To achieve high performance, the dynamics model must accurately predict the rewards ahead for multiple time steps. They approach this using a latent dynamics model for its fast querying capabilities. Moreover, they propose a multi-step variational inference objective named latent overshooting.

PlaNet learns a transition model  $p(s_t|s_{t-1}, a_{t-1})$ , observation model  $p(o_t|s_t)$ , and reward model  $p(r_t|s_t)$  from previously experienced episodes. All of them are Gaussians parameterised by neural networks. The observation model provides a rich training signal but is not used for planning. It also learns an encoder  $q(s_t|o_{\leq t}, a_{<t})$  to infer an approximate belief over the current hidden state from the history using filtering. This model can be thought of as a non-linear Kalman filter or sequential VAE. The transition model consists of both stochastic and deterministic paths as experiments show that these are crucial for successful planning.

Despite its generality, the purely stochastic transitions make it difficult for the transition model to reliably remember information for multiple time steps. In theory, this model could learn to set the variance to zero for some state components, but the optimization procedure may not find this solution. This motivates including a deterministic sequence of activation vectors  $\{h_t\}_{t=1}^T$  that allow the model to access not just the last state but all previous states deterministically. The authors use

such a model, shown in fig. 34c, that they name recurrent state-space model (RSSM). Intuitively, one can understand this model as splitting the state into a stochastic part  $s_t$  and a deterministic part  $h_t$ , which depend on the stochastic and deterministic parts at the previous time step through the recurrent neural network (RNN).

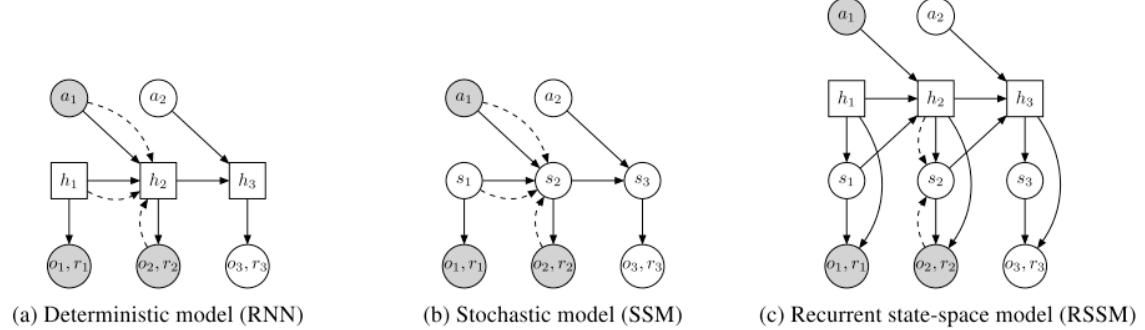


Fig. 34. Latent dynamics model designs [20]. In this example, the model observes the first two time steps and predicts the third. Circles represent stochastic variables and squares deterministic variables. Solid lines denote the generative process and dashed lines the inference model. (a) Transitions in a recurrent neural network are purely deterministic. This prevents the model from capturing multiple futures and makes it easy for the planner to exploit inaccuracies. (b) Transitions in a state-space model are purely stochastic. This makes it difficult to remember information over multiple time steps. (c) This model splits the state into stochastic and deterministic parts, allowing the model to robustly learn to predict multiple futures.

Given these components, the authors implement the policy as a planning algorithm that searches for the best sequence of future actions. In contrast to model-free and hybrid reinforcement learning algorithms, the authors do not use a policy or value network.

The cross entropy method [3] (CEM) is used to search for the best action sequence under the model. The authors decided on this algorithm because of its robustness and because it solved all considered tasks when given the true dynamics for planning. CEM is a population-based optimization algorithm that infers a distribution over action sequence that maximize the objective. Starting from zero mean and unit variance time-dependent diagonal Gaussian belief over optimal action sequence up to  $H$ , the length of the planning horizon, it repeatedly sample  $J$  candidate action sequences, evaluate them under the model, and re-fit the belief, mean and variance, to the top best  $K$  action sequences. After  $I$  iterations, the planner returns the mean of the belief for the current time step. Importantly, after receiving the next observation, the belief over action sequences starts from zero mean and unit variance again to avoid local optima.

To evaluate a candidate action sequence under the learned model, PlaNet samples a state trajectory starting from the current state belief and sum the mean rewards predicted along the sequence. Since CEM is a population-based optimizer, the authors found it sufficient to consider a single trajectory per action sequence and thus focus the computational budget on evaluating a larger number of different sequences. Because the reward is modeled as a function of the latent state, the planner can operate purely in latent space without generating images, which allows for fast evaluation of large batches of action sequences.

Since the agent may not initially visit all parts of the environment, it needs to iteratively collect new experience and refine the dynamics model. It does so by planning with the partially

trained model. Starting from a small amount of  $S$  seed episodes collected under random actions, the authors train the model and add one additional episode to the data set every  $C$  update steps.

Typical variational bound for learning and inference in latent sequence models, as shown in fig. 35a, contains reconstruction terms for the observations and KL-divergence regularisers for the approximate posteriors. A limitation of this objective is that the stochastic path of the transition function  $p(s_t|s_{t-1}, a_{t-1})$  is only trained via the KL-divergence regularisers for one-step predictions: the gradient flows through  $p(s_t|s_{t-1}, a_{t-1})$  directly into  $q(s_{t-1})$  but never traverses a chain of multiple  $p(s_t|s_{t-1}, a_{t-1})$ . If one could train a model to make perfect one-step predictions, it would also make perfect multi-step predictions, so this would not be a problem. However, when using a model with limited capacity and restricted distributional family, training the model only on one-step predictions until convergence does in general not coincide with the model that is best at multi-step predictions. For successful planning, accurate multi-step predictions are needed.

The authors generalize the standard variational bound to latent overshooting, as shown in fig. 35c, which trains all multi-step predictions in latent space. Using only terms in latent space results in a fast regulariser that can improve long-term predictions, it encourages consistency between one-step and multi-step predictions, and is compatible with any latent sequence model.

The authors found that several dynamics models benefit from latent overshooting, although their final agent using the RSSM model does not require it.

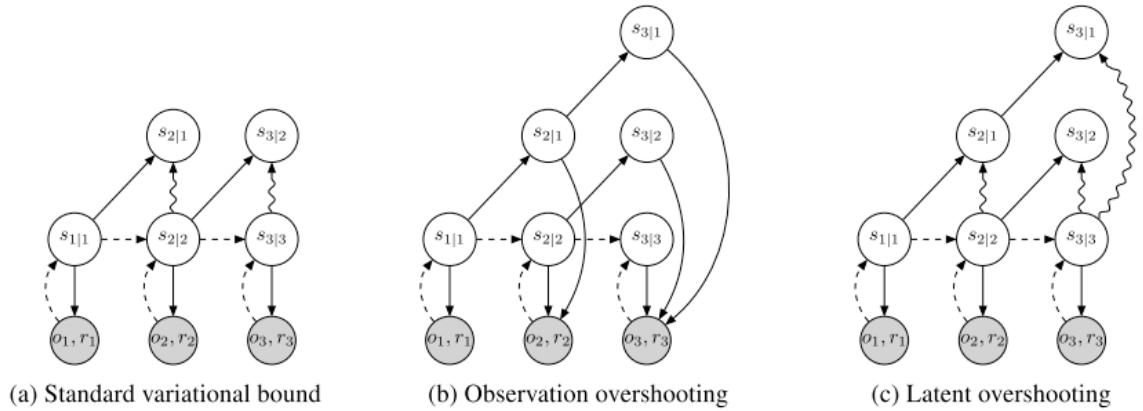


Fig. 35. Unrolling schemes [20]. The labels  $s_{i|j}$  are short for the state at time  $i$  conditioned on observations up to time  $j$ . Arrows pointing at shaded circles indicate log-likelihood loss terms. Wavy arrows indicate KL-divergence loss terms. (a) The standard variational objectives decodes the posterior at every step to compute the reconstruction loss. It also places a KL on the prior and posterior at every step, which trains the transition function for one-step predictions. (b) Observation overshooting decodes all multi-step predictions to apply additional reconstruction losses. This is typically too expensive in image domains. (c) Latent overshooting predicts all multi-step priors. These state beliefs are trained towards their corresponding posteriors in latent space to encourage accurate multi-step predictions.

Using only pixel observations, PlaNet’s agent solves continuous control tasks from DeepMind control suite (see fig. 36) with contact dynamics, partial observability, and sparse rewards, which exceed the difficulty of tasks that were previously solved by planning with learned models. PlaNet uses substantially fewer episodes and reaches final performance close to and sometimes higher than strong model-free algorithms.

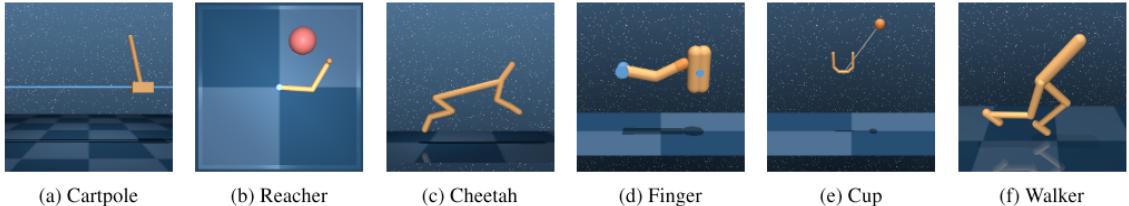


Fig. 36. DeepMind Control Suite: image-based control domains used in PlaNet’s experiments [20].

Within 100 episodes, PlaNet outperforms the policy-gradient method A3C [36] trained from proprioceptive states for 100,000 episodes, on all tasks. After 500 episodes, it achieves performance similar to D4PG [1], trained from images for 100,000 episodes, except for the finger task. PlaNet surpasses the final performance of D4PG with a relative improvement of 26% on the cheetah running task.

Moreover, the authors trained a single agent on all six tasks. The agent is not told which task it is facing, it needs to infer this from the image observations. The agent solves all tasks while learning slower compared to individually trained agents. This indicates that the model can learn to predict multiple domains, regardless of the conceptually different visuals.

PlaNet is a working example of a model-based agent that learns a latent dynamics model from high-dimensional image observations and chooses actions by fast planning in latent space. The authors show that their agent succeeds at several continuous control tasks from image observations, reaching performance that is comparable to the best model-free algorithms while using 200 times fewer episodes and similar or less computation time. The results show that learning latent dynamics models for planning in image domains is a promising approach. This thesis will adopt PlaNet to its objectives.

### 3.2.3. Model-Based Reinforcement Learning for Atari

In this paper [29], the authors explore how video prediction models can enable RL agents to solve Atari games with orders of magnitude fewer interactions than model-free methods. They describe Simulated Policy Learning (SimPLe), a complete model-based deep RL algorithm based on video prediction models, and present a comparison of several model architectures, including a novel architecture that yields the best results in Atari setting.

SimPLe, apart from an Atari 2600 emulator environment  $env$ , use a neural network simulated environment  $env'$  which they call a world model. The authors find out that crucial decisions in the design of world models are: inclusion of stochasticity, clipping reconstruction loss to a constant and, because simulator  $env'$  is supposed to consume its own predictions from previous steps which will be imperfect due to compounding errors, to mitigate a problem of the model drifting out of the area of its applicability they randomly replace in training some frames of the input by the model prediction from the previous step.

The environment  $env'$  shares the action space and reward space with  $env$  and produces visual observations in the same format, as it will be trained to mimic  $env$ . The authors principal

aim is to train a policy  $\pi$  using a simulated environment  $env'$  so that  $\pi$  achieves good performance in the original environment  $env$ . Using short rollouts for a policy training is crucial to mitigate the compounding errors under the model. To ensure exploration SimPLe starts training rollouts from randomly selected states taken from the real data buffer.

In this training process the authors aim to use as few interactions with  $env$  as possible. The initial data to train  $env'$  comes from random rollouts of  $env$ . As this is unlikely to capture all aspects of the environment, they use the data-aggregation iterative method.

Experiments evaluate SimPLe on a range of Atari games and show that it achieves competitive results compared to model-free baselines with only 100K interactions between the agent and the environment, which corresponds to about two hours of real-time play. SimPLe is significantly more sample-efficient than a highly tuned version of the state-of-the-art Rainbow algorithm [22] on almost all games. In particular, in low data regime of 100k samples, on more than half of the games, SimPLe achieves a score which Rainbow requires at least twice as many samples. In the best case of Freeway, it is more than 10x more sample-efficient.

SimPLe is a similar approach to model-based RL like in World Models. The authors train the dynamics model and use it to generate new experience, the same as in World Models, but in observations space, opposite to World Models.

Although this thesis share the goal of sample efficient RL via model-based planning and learning in complex environments, like Atari games, with SimPLe, it tries to accomplish it in fundamentally different way. The thesis' solution focuses on train accurate transitions model in the latent state-space to enable fast simulation and search using this model. Pixel-perfect reconstructions are not a concern.

### 3.3. Monte-Carlo planning

Planning is a natural and powerful approach to decision making problems with known dynamics. For instance, Monte-Carlo Tree Search methods [32] have been used for complex search problems, such as the game of Go [50]. Moreover, planning carries the promise of increasing performance just by increasing the computational budget for searching for actions. The first paper below describes current state of the art search method and the other one explores the idea of Monte-Carlo planning with a learned model.

#### 3.3.1. A general reinforcement learning algorithm that masters chess, shogi, and Go

AlphaZero [49] replaces the handcrafted knowledge and domain-specific augmentations used in traditional game-playing programs with deep neural networks, a general-purpose reinforcement learning algorithm, and a general-purpose tree search algorithm. Instead of handcrafted move ordering and evaluation function heuristics, AlphaZero infers policy and state value with a deep neural network. This neural network takes the board position as an input and outputs a vector of move probabilities for each action and a scalar value estimating the expected outcome of

the game. AlphaZero learns these move probabilities and value estimates entirely from self-play. These are then used to guide its search in future games.

Instead of an alpha-beta search with domain-specific enhancements, as common in board games playing programs, AlphaZero uses a general-purpose Monte Carlo tree search [32] algorithm. Each search consists of a series of simulated games that traverse a tree from a root state until a leaf state is reached. Each simulation proceeds by selecting in each state a move with low visit count (not previously frequently explored), high move probability and high value (averaged over the leaf states of simulations that selected this move previously) according to the current neural network. The search returns a vector representing a probability distribution over moves from the root state.

The parameters of the deep neural network in AlphaZero are trained by reinforcement learning from self-play games, starting from randomly initialized parameters. Each game is played by running the search, described above, from the current position and then selecting a move either proportionally (for exploration) or greedily (for exploitation) according to returned move probabilities, which are normalised visit counts at the root state. At the end of the game, the terminal position is scored according to the rules of the game to compute the game outcome: -1 for a loss, 0 for a draw, and +1 for a win. The neural network parameters are updated to minimize the error between the predicted outcome and the game outcome, and to maximize the similarity of the predicted move probabilities to the search probabilities. The updated parameters are used in subsequent games of self-play.

Starting from random play, AlphaZero convincingly defeats a world champion programs in the games of chess and shogi as well as Go. Especially, the game of chess represents the pinnacle of artificial intelligence research over several decades. State-of-the-art programs are based on powerful engines that search many millions of positions, leveraging handcrafted domain expertise and sophisticated domain adaptations. AlphaZero is a generic reinforcement learning and search algorithm, originally devised for the game of Go, that achieve superior results within a few hours, searching 1/1000 as many positions and it is given no domain knowledge except the rules of chess. Furthermore, the same algorithm was applied without modification to the more challenging game of shogi, again outperforming state-of-the-art programs within a few hours.

AlphaZero is the state-of-the-art general search algorithm, but it requires a dynamics model of an environment in order to work. The model is not always available, but could be learned in such cases. This thesis tries to enable AlphaZero to work with the learned model. This work is described in the next chapters.

### 3.3.2. Value Prediction Network

This paper [41] proposes a novel deep reinforcement learning architecture, called Value Prediction Network (VPN), which integrates model-based (i.e. learning an abstract dynamics model) and model-free (i.e. mapping the learned abstract states to rewards and values) RL methods into a single neural network. In contrast to typical model-based RL methods, VPN learns the

dynamics model of an abstract state space sufficient for predicting future rewards and values, rather than future observations.

In order to train VPN, the authors propose a combination of temporal-difference search [48] and n-step Q-learning [36]. In brief, VPN learns to predict values via Q-learning and rewards via supervised learning. At the same time, VPN perform lookahead planning to choose actions and compute bootstrapped target Q-values.

VPN has the ability to simulate the future and plan based on the simulated future abstract-states. Although many existing planning methods (e.g. MCTS) can be applied to the VPN, the authors implement a simple planning method which performs Monte-Carlo rollouts using the VPN up to a certain depth and aggregates all intermediate value estimates as described in the paper [41]. The planning procedure estimates the current state Q-values which an agent uses to decide on the next action.

Experimental results show that VPN has several advantages over both model-free and model-based baselines in a stochastic navigation task where careful planning is required but building an accurate observation-prediction model is difficult. Furthermore, VPN outperforms Deep Q-Network [37], strong model-free baseline method, on several Atari games even with short-lookahead planning, demonstrating its potential as a new way of learning a good state representation.

VPN uses an abstract model of rewards to augment model-free learning with good results on a number of Atari games. However, this method does not actually aim to model or predict future environment's states and achieves clear but relatively modest gains in sample-efficiency. On the other hand, it is an example of simple tree search algorithm which can successfully use the learned model for planning, something that this thesis could base on.

## 4. PLANNING WITH LEARNED MODEL

In this section three architectures are described. All of them involve a similar model learning approach, but differ substantially in technical details and planning algorithms. All architectures use computationally efficient latent space environment models that make predictions at a higher level of spatial abstraction, than at the level of raw pixel observations. Such models substantially reduce the amount of computation required to make predictions, as future states can be represented much more compactly. In order to increase model accuracy and robustness, all models explicitly model uncertainty in state transitions using stochastic nodes [5]. The goal stays the same: train a sufficiently accurate latent space environment's dynamics model, or such that accurately predicts future latent states and rewards to predefined cut-off point in time, and use it to plan and solve the environment. Those architectures, and experiments carried out on them, present evolution of ideas towards this goal.

Before all of that, the code architecture and the framework that was created to accelerate this research are described.

### 4.1. *HumbleRL* framework

Reinforcement Learning scientists tend to write the entire code from scratch by themselves, instead of using existing RL frameworks. This is justified by the fact, that the commonly available frameworks are not flexible enough for intended experiments or require a specific backend like TensorFlow, which might be disfavored. HumbleRL [27] was created with this problem in mind. Its simple API allows to perform a variety of RL experiments without any restrictions on the algorithms used. Since the backend is not tied to any specific technology, it is possible to mix different neural network frameworks or not use them at all. HumbleRL provides the boilerplate code of RL loop in fig.21 and determines the common interfaces between an agent and an environment, the rest is designed by the user.

#### 4.1.1. *Architecture*

Framework architecture is depicted in fig. 41. An agent is represented by the Mind class. The Mind encapsulates action planning logic and provides it via the plan method. In order to learn, the agent acts in the world represented by the Environment class. The Environment provides methods for resetting, taking steps, rendering and getting information about the world. The framework includes a factory function that creates and returns e.g. wrapped OpenAI Gym environment. The agent is not usually presented with raw environment observations. Instead, it looks at states preprocessed by the Interpreter. Different interpreters can be joined together with the ChainInterpreter class. It acts as a preprocessing pipeline, with each subsequent interpreter using the output of a previous one as an input.

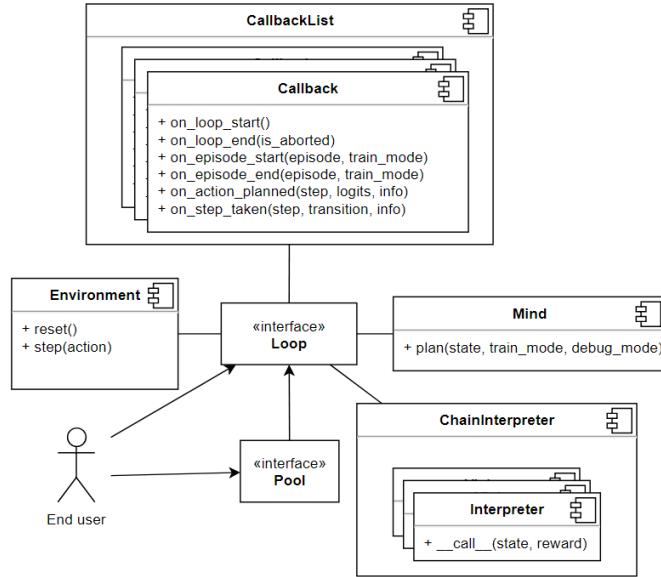


Fig. 41. HumbleRL architecture

Framework user does not need to call all of those methods directly, those are utilized by the loop function. This function gets an action from the Mind, executes it in the Environment and then next observation is preprocessed with the Interpreter in preparation for the next step. To extend basic loop functionality, user can define callbacks that implement the `Callback` interface. Callbacks can react to events:

- at the beginning and ending of the loop,
- at the beginning and ending of each episode,
- after action is planned by the Mind,
- after step is taken in the Environment.

Callbacks are accumulated in the `CallbackList`. The entire loop function logic is shown in fig. 42. Parallel version of loop function is available as the `pool` function. It uses predefined number of workers to execute a pool of Minds in their own Environments in parallel.

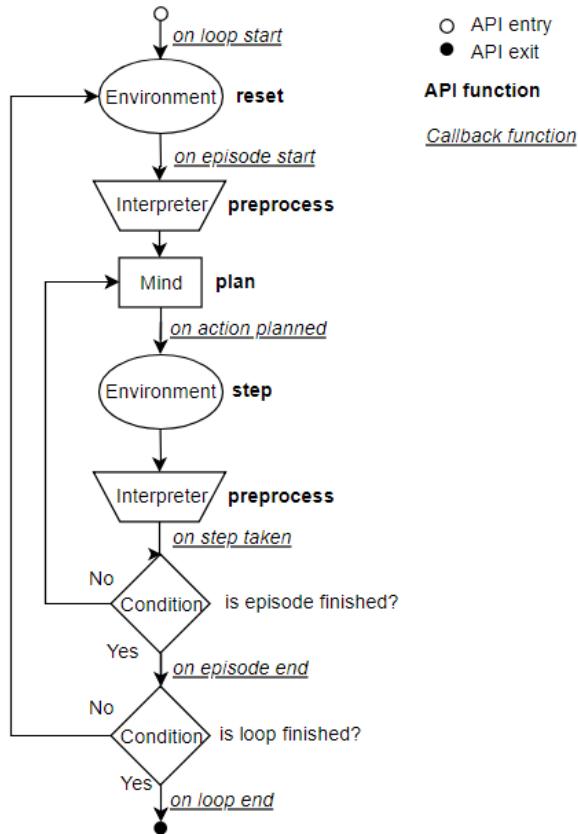


Fig. 42. HumbleRL loop function overview

World Models and AlphaZero implementations use this framework and can be joined together into one architecture.

#### 4.2. Original World Models

This section describes the original World Models [17] algorithm which uses a learned model to create abstract task representation which, in turn, is used to solve the task by a lightweight controller. This architecture will be coupled with the AlphaZero which uses the learned model to simulate experience in the next section.

##### 4.2.1. World model

A simple model inspired by human cognitive system is used. In this model, an agent has a visual sensory module that compresses observations into a small representative code. It also has a memory module that makes predictions about future codes based on historical information. Finally, the agent has a decision-making component, called the controller module, that decides what actions to take based only on the representations created by its vision and memory modules. This architecture allows for training of a large neural network to tackle RL tasks by dividing the agent into a large world model and a small controller model. First a large neural network learns to model the agent's world in an unsupervised manner, and only then the smaller controller focuses

on the credit assignment problem on a smaller search space of controller's parameters, while not sacrificing capacity and expressiveness via the larger world model.

An environment provides the agent with a high dimensional input observation at each time step. It is a 2D image frame that is part of a video sequence. The vision module role, as already mentioned, is to learn an abstract representation of each observed input frame. It uses a simple Variational Autoencoder [28] and, as described in details in theoretical background chapter, is trained to encode each frame into low dimensional latent vector by minimizing the difference between a given frame and the reconstructed version of the frame produced by the decoder.

The Memory module purpose is to compress the information what happens over time in its hidden state and enable simulation of the environment. To do this, it is trained to model environment's dynamics, predicting a future latent state from history of previous latent states, as a mixture of Gaussians. It models latent states with probability distribution to model uncertainty in the environment, but also create more robust environment's representation [5]. Uncertainty can originate not only from fundamental stochastic nature of the environment, but also partial observability.

The model is implemented as a recurrent neural network with the Mixture Density Network (MDN) on top of a RNN's hidden state. In literature this architecture is called MDN-RNN [14]. It will be further referred to as the stochastic Memory module.

Figure 43 depicts the world model, the Vision and Memory modules interconnection, in graphical form. More specifically, the world model components are:

- Deterministic hidden state model:  $h_t = f(h_{t-1}, z_t, a_t)$
- Stochastic latent state model:  $z_{t+1} \sim p(z_{t+1}|h_t) = \sum_c \pi_c(h_t)p(z_{t+1}|h_t, c)$
- Observation model (decoder):  $o_t \sim p(o_t|z_t)$
- Approximate state posterior (encoder):  $z_t \sim q(z_t|o_t)$

where  $o$ ,  $z$  and  $a$  are high-dimensional observations, latent states and actions respectively.  $f(h_{t-1}, z_t, a_t)$ , the hidden state model, is implemented as a recurrent neural network and  $h_t$  is its hidden state. The latent state model is a mixture of Gaussians with mean and variance parameterised by a feed-forward neural network.  $c$  is a mixture's component and  $\pi(h)$  is a normalized vector of mixing coefficients as a function of the RNN's hidden state. The observation model is Bernoulli distribution parameterised by a deconvolutional neural network. Since the model is non-linear, directly computing the state posteriors is intractable. Instead, an encoder  $q$  is used to infer approximate state posteriors from past observations and actions, where  $q(s_t|h_t, o_t)$  is a diagonal Gaussian with mean and variance parameterised by a convolutional neural network followed by a feed-forward neural network.

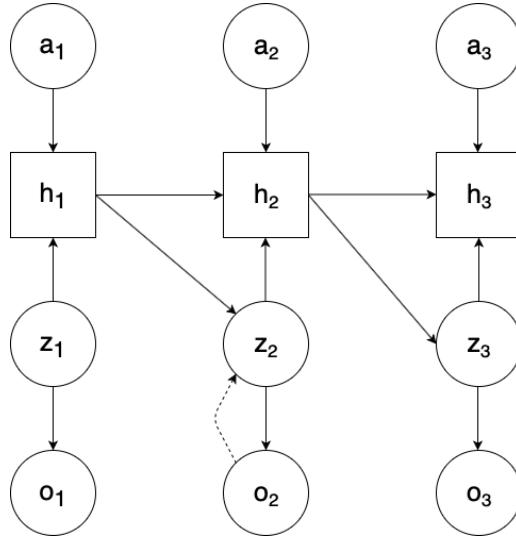


Fig. 43. World Models probabilistic graphical model: solid arrows describe the predictive model, dotted arrow describes the inference model, stochastic nodes are circles and squares depict deterministic nodes.

#### 4.2.2. Controller

The Controller module is responsible for determining the course of actions to take in order to maximize the expected cumulative reward of the agent during an episode in the environment. Because it bases its decisions on abstract world representation, learnt by the Vision and Memory modules, it can be deliberately made as simple and small as possible and trained separately from Vision and Memory modules, so that most of the agent's complexity resides in the world model. The Controller module is a simple single layer linear model that maps a features vector constructed from latent state of the Vision module  $z$  and hidden state of the Memory module  $h$  directly to action  $a$  at each time step:

$$a_t = W_c[z_t \ h_t] + b_c$$

where  $W_c$  and  $b_c$  are Controller's weights matrix and biases vector that maps the concatenated features vector  $[z_t \ h_t]$  to the output action vector  $a_t$ . The action vector is  $|A|$ -dimensional, where  $|A|$  is number of legal actions in an environment. It encodes predicted score of each action in the state and is used to decide which action to choose in the environment. Maximum action is taken, which means that greedy policy is used. It is true for training and testing phases. Because evolutional strategy algorithm is used for training, the fact that is described below, it doesn't impair exploration which is done on parameters level.

#### 4.2.3. Data collection

To train Vision and Memory modules first collection of 10,000 random rollouts of the environment are gathered to create a dataset. An agent is acting randomly to explore the environment multiple times and records the random actions taken and the resulting observations from the en-

vironment. This dataset is used to train the Vision module. Next, it is used to process each frame into its latent state to prepare a dataset for the Memory module training, which works entirely in the latent space.

The Controller module is trained using evolutional strategy, described below in the implementation details section, which rollouts the environment in each epoch to evaluate population.

#### 4.2.4. *Preprocessing*

Each frame, before it is used for any training, is central cropped if a frame from an environment includes some kind of border which doesn't inform an agent in anyway. This operation depends on a specific environment. It is then resized to 64 x 64 pixels for all environments. All three colour channels are preserved. Actions are one-hot encoded and default 4 action repeat from OpenAI Gym [4] is used as common in reinforcement learning [37] to reduce the planning horizon and provide a clearer learning signal to the model.

#### 4.2.5. *Implementation details*

HumbleRL is used to implement the original World Models architecture from the paper. This allows for easy adjustments for experiments purposes and to couple the world model with AlphaZero implementation in HumbleRL. An agent exploring an environment (Mind) and a call-back are used to gather transitions and save them to an external storage. The framework allows to focus strictly on collecting trajectories and not worry about agent-environment interactions.

Collected transitions are used to train the Vision and Memory components. The popular LSTM architecture [24] implements the Memory module RNN with 256 hidden units. The MDN is composed of 5 16-dimensional Gaussians with mean and log standard deviation parameterised by one layer feed-forward neural networks with linear activations. The Vision module neural networks are convolutional and deconvolutional neural networks shown in figure 44 where latent state size,  $\mu$  and  $\sigma$  vectors dimensionality, equals 16.

For Vision training Keras [8] framework is used to adjust the parameters by the stochastic gradient descent on a standard Evidence Lower Bound loss. For Memory training PyTorch [42] framework is used, since it is easier to work with recurrent neural networks than in Keras, to, again, adjust the parameters by the stochastic gradient descent on a negative log-probability loss of the mixture density network. HumbleRL is not constricted to work with any particular deep learning library, so it is not a problem to mix the solutions, as long as trained models are wrapped in HumbleRL's interfaces. The Vision module is trained using the Adam optimizer [30] with a learning rate of  $10^{-3}$  and  $\epsilon = 10^{-7}$  on batches of 256 images. The Memory module is trained using the Adam optimizer [30] too with a learning rate of  $10^{-3}$  and  $\epsilon = 10^{-8}$  on batches of 128 sequence chunks of length 1000.

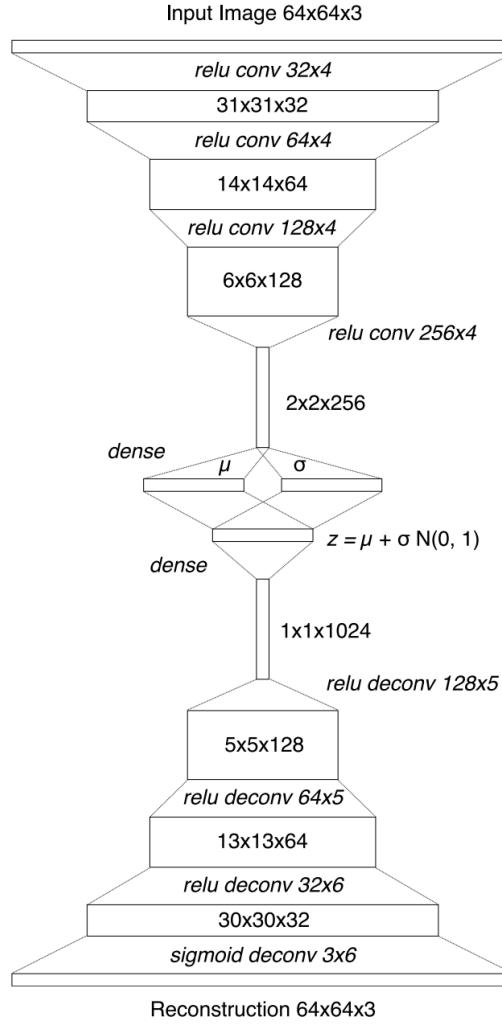


Fig. 44. World Models VAE neural network architecture [17]

The Controller module gets a feature vector as its input, consisting of latent state  $z$  and the hidden state of the MDN-RNN. In all environments, this hidden state is the output vector  $h$  of the LSTM.

Covariance-Matrix Adaptation Evolution Strategy (CMA-ES) [21] is used for the Controller module training. It evolves the weights  $W_c$  and biases  $b_c$  of the module. A population size of 64 is used and each agent plays in the environment 5 episodes. The fitness value for the agent is the average cumulative reward of the 5 episodes.

Hiper-parameters presented are used as defaults in experiments described in the next chapter.

#### 4.3. World Models and AlphaZero

World Models' agent [17] successfully plans using a learned model where the model is used to generate simulated experience on which the policy is trained. This section describes attempt to adjust and utilize the world model part of the agent in the AlphaZero search algorithm. This is different application of the model than in the original paper, where only future latent states

and done flag are predicted, and therefore the world model needs to be extend with a reward predictor and the controller is replaced by AlphaZero.

#### 4.3.1. World model

The world model part, Vision and Memory modules, architecture stays the same with a slight change. Because benchmarks include only deterministic environments and AlphaZero in its original form can only work with a deterministic dynamics model, this architecture use the Memory module without the MDN and with a linear model instead. In fact, it uses two linear models to output the next latent state and reward.

- Deterministic hidden state model:  $h_t = f(h_{t-1}, s_t, a_t)$
- Deterministic latent state model:  $z_{t+1} = f(h_t)$
- Deterministic reward model:  $r_{t+1} = f(h_t)$
- Observation model (decoder):  $o_t \sim p(o_t | z_t)$
- Approximate state posterior (encoder):  $z_t \sim q(z_t | o_t)$

where  $f$  in the latent state model and the reward model are the linear models. It will be further referred to as the deterministic Memory module.

#### 4.3.2. Controller

AlphaZero [49] is very similar to the MCTS algorithm, explained in the previous chapter. The selection and evaluation phases are modified though. AlphaZero uses policy and value networks to guide its search. Each edge in the search tree,  $(s, a, s')$  where  $s$  is a state,  $a$  is an action and  $s'$  is the next state, stores a prior probability of choosing it  $P(s, a)$ , a visit count  $N(s, a)$ , an action-value  $Q(s, a)$  and a reward  $R(s, a, s')$  of transition represented by this edge. Each simulation starts from the root state,  $s_t$  at depth or time step  $t = 0$ , and iteratively selects moves that maximise an upper confidence bound  $Q(s, a) + U(s, a)$ , where  $U(s, a) \propto P(s, a)/(1 + N(s, a))$  [50], until a leaf node is encountered. This leaf position is expanded by the world model to generate both the next state and reward,  $s_{t'}$  and  $r_{t'}$  at depth or time step  $t'$ , and evaluated by the networks to generate both prior actions probabilities and its value,  $P(s_{t'}, \cdot)$  and  $V(s_{t'})$ . Each edge traversed in the simulation is updated to increment its visit count  $N(s_t, a_t)$ , and to update its action-value to the mean evaluation over these simulations:

$$Q(s_t, a_t) = 1/N(s_t, a_t) \sum_{t'} \left[ \sum_{i=t+1}^{t'} r_i + V(s_{t'}) \right]$$

Figure 45 depicts this process. After multiple simulations, which could be stopped after the absolute number of simulations or after timeout, the result is a vector of search probabilities recommending moves to play,  $\pi$ , proportional to the visit count for each move,  $\pi_a \propto N(s, a)$ . How these are used to choose an action to take by the agent is discussed in the next section.

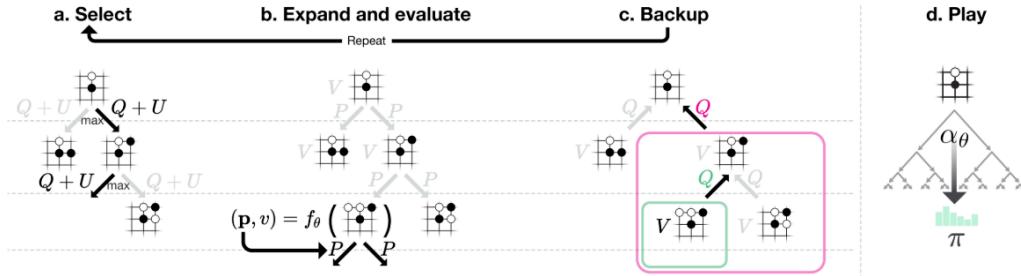


Fig. 45. Monte-Carlo tree search in AlphaZero [50]

#### 4.3.3. Implementation details, data collection and preprocessing

The world model implementation details, preprocessing and data collection for Vision and Memory modules are the same as in the original World Models. The world model latent state size is 16 and the LSTM hidden size is 256. Input to value and policy networks is the same concatenated features vector of latent and hidden states,  $[z_t \ h_t]$ .

The AlphaZero controller uses the world model for simulations. The Vision module is used as the Interpreter which encodes incoming observations into latent space. The Memory module, wrapped in the MDP interface from HumbleRL, is used in the expansion phase of AlphaZero. The Mind class, which is implemented by the AlphaZero algorithm, returns actions' scores. These are actions visit counts from the root state node, which are then used to choose an action by a policy. During training actions are sampled with probability proportional to these visit counts for 12 warm-up steps at the beginning of each episode to enhance exploration and after warm-up a greedy policy is used. During testing always greedy policy is used, which picks an action which was visited most often. Pseudo-code written in Python of the search algorithm in the Mind is shown below:

```

1 def plan(self, state):
2     # Get/create root node
3     root = self.query_tree(state)
4
5     # Perform simulations
6     simulations = 0
7     start_time = time()
8     while time() < start_time + self.timeout and simulations < self.max_simulations:
9         # Simulate
10        simulations += 1
11        leaf, path = self.simulate(root)
12
13        # Expand and evaluate
14        value = self.evaluate(leaf)
15
16        # Backup value
17        self.backup(path, value)
18
19        # Get actions' visit counts
20        actions = np.zeros(self.model.action_space.num)
21        for action, edge in root.edges.items():
22            actions[action] = edge.num_visits
23
24    return actions

```

AlphaZero value and policy networks are two linear models trained separately from the world model from games played by the agent, called self-play. In each position  $s$  during self-play,

an MCTS search is executed, guided by the policy network. The MCTS search at the end outputs probabilities,  $\pi$ , of playing each move. These search probabilities usually select much stronger moves than the raw move probabilities of the policy network. The MCTS search may therefore be viewed as a powerful policy improvement operator [52]. Self-play with search – using the improved MCTS-based policy to select each move by the agent, then using the game cumulative reward, or the return, as a sample of the value at each time step – may be viewed as a powerful policy evaluation operator. The main idea of this reinforcement learning algorithm is to use these search operators repeatedly in a policy iteration framework [50]. The linear models' parameters are updated to make the move probabilities and values more closely match the improved search probabilities and self-play cumulative rewards at each step. These new parameters are used in the next iteration of self-play to make the search even stronger.

The agent's experience and score statistics used for training are gathered using callbacks during the self-play phase. Maximum of 1000 latest games are kept. The models training phase takes place after 10 self-play games and lasts for 5 epochs. The training phase is performed using the Keras [8] framework. Specifically, the parameters are adjusted by the stochastic gradient descent on a loss function that sums over mean squared errors of the value network, cross-entropy losses of the policy network and L2 weights regularization scaled by a factor of  $10^{-4}$  in batches of 256 examples, where each example is a features vector of a state, a value sample from self-play and a MCTS action probabilities vector. The Nesterov's momentum optimizer [51] is used with a learning rate of  $10^{-2}$  and a momentum of 0.9.

Hiper-parameters presented are used as defaults in experiments described in the next chapter.

#### 4.4. PlaNet

PlaNet (Deep Planning Network) [20] shows working example of a planning agent that searches for the best sequence of future actions using a learned model in continuous control tasks. This is close to what this work tries to accomplish, but for a different type of environments. This section describe how it was utilized in episodic discrete tasks.

##### 4.4.1. World model

This architecture uses recurrent state space model (RSSM) which is similar to what World Models does. This latent dynamics model is designed with both deterministic and stochastic components [5]. Original experiments indicate having both components to be crucial for high planning performance. It also uses a generalized variational bound that include multi-step predictions. Using only terms in latent space results in a fast regularizer that can improve long-term predictions [20].

The same as in World Models, the model is provided with image observations. It even uses the same Variational Autoencoder with the same neural network architecture to encode the

observations into latent space. The difference lies in the dynamics model shown in figure 46 with following components:

- Deterministic hidden state model:  $h_t = f(h_{t-1}, s_{t-1}, a_{t-1})$
- Stochastic latent state model:  $s_t \sim p(s_t|h_t)$
- Observation model (decoder):  $o_t \sim p(o_t|h_t, s_t)$
- Reward model:  $r_t \sim p(r_t|h_t, s_t)$
- Approximate state posterior (encoder):  $s_t \sim q(s_t|o_{\leq t}, a_{<t}) = \prod_{i=1}^t q(s_i|h_i, o_i)$

where  $o$ ,  $s$  and  $a$  are high-dimensional observations, latent states and actions respectively.  $f(h_{t-1}, s_{t-1}, a_{t-1})$ , the deterministic state model, is implemented as a recurrent neural network and  $h_t$  is its hidden state. The latent state model is Gaussian with mean and variance parameterised by a feed-forward neural network, the observation model is Gaussian with mean parameterised by a deconvolutional neural network and identity covariance, and the reward model is a scalar Gaussian with mean parameterised by a feed-forward neural network and unit variance. Since the model is non-linear, directly computing the state posteriors is intractable. Instead, an encoder  $q$  is used to infer approximate state posteriors from past observations and actions, where  $q(s_t|h_t, o_t)$  is a diagonal Gaussian with mean and variance parameterised by a convolutional neural network followed by a feed-forward neural network.

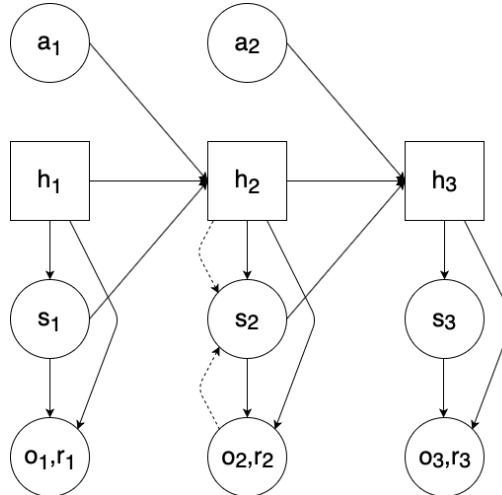


Fig. 46. PlaNet probabilistic graphical model: solid arrows describe the predictive model, dotted arrow describes the inference model, stochastic nodes are circles and squares depict deterministic nodes.

There are two deviations from the World Models architectures:

- Hidden states are “shifted” to the right, so the previous state and action are used to predict the current hidden state  $h$  and it is then used to predict the current latent state  $s$ .
- VAE’s decoder and encoder (named Vision in World Models) use the dynamics model’s hidden state for prediction and inference. This way the likelihood and the posterior are conditioned on past observations too, as opposed to World Models.

Intuitively, this model can be understood as splitting the state into a stochastic part and a deterministic part, which depend on the stochastic and deterministic parts at the previous time

step through the RNN. Importantly, all information about the observations must pass through the sampling step of the encoder to avoid a deterministic shortcut from inputs to reconstructions.

#### 4.4.2. Planner

The agent plans using the cross entropy method (CEM) [3] to search for the best action sequence under the model. CEM is a population-based optimization algorithm that infers a distribution over action sequences that maximize the objective. First, a time-dependent diagonal Gaussian belief over optimal action sequences gets initialized:  $a_{t:t+H} \sim \text{Normal}(\mu_{t:t+H}, \sigma_{t:t+H}^2 I)$ , where  $t$  is the current time step of the agent and  $H$  is the length of the planning horizon. Starting from zero mean and unit variance, it is used to sample candidate action sequences. To evaluate a candidate action sequence under the learned model a trajectory starting from the current state is sampled by the model using the action sequence as an input and the predicted rewards are summed along the sequence. This sum is used as the action sequence fitness score. Since it is a population-based optimizer, it is sufficient to consider a single trajectory per action sequence and thus focus the computational budget on evaluating a larger number of different sequences. Then, the belief gets re-fitted to the best sequences and next optimization iteration starts the same way. At the end, the planner returns the mean of the belief for the current time step  $\mu_t$  which is then used by the policy to choose discrete action. Importantly, after receiving the next observation, the belief over action sequences starts from zero mean and unit variance again to avoid local optima. Because the reward is modeled as a function of the latent state, the planner can operate purely in latent space without generating images, which allows for fast evaluation of large batches of action sequences. The algorithm is presented in the figure 47 below.

|               |                                     |                               |                      |
|---------------|-------------------------------------|-------------------------------|----------------------|
| <b>Input:</b> | $H$ Planning horizon distance       | $q(s_t   o_{\leq t}, a_{<t})$ | Current state belief |
|               | $I$ Optimization iterations         | $p(s_t   s_{t-1}, a_{t-1})$   | Transition model     |
|               | $J$ Candidates per iteration        | $p(r_t   s_t)$                | Reward model         |
|               | $K$ Number of top candidates to fit |                               |                      |

```

Initialize factorized belief over action sequences  $q(a_{t:t+H}) \leftarrow \text{Normal}(0, \mathbb{I})$ .
for optimization iteration  $i = 1..I$  do
    // Evaluate  $J$  action sequences from the current belief.
    for candidate action sequence  $j = 1..J$  do
         $a_{t:t+H}^{(j)} \sim q(a_{t:t+H})$ 
         $s_{t:t+H+1}^{(j)} \sim q(s_t | o_{1:t}, a_{1:t-1}) \prod_{\tau=t+1}^{t+H+1} p(s_\tau | s_{\tau-1}, a_{\tau-1}^{(j)})$ 
         $R^{(j)} = \sum_{\tau=t+1}^{t+H+1} \mathbb{E}[p(r_\tau | s_\tau^{(j)})]$ 
    // Re-fit belief to the  $K$  best action sequences.
     $\mathcal{K} \leftarrow \text{argsort}(\{R^{(j)}\}_{j=1}^J)_{1:K}$ 
     $\mu_{t:t+H} = \frac{1}{K} \sum_{k \in \mathcal{K}} a_{t:t+H}^{(k)}$ ,  $\sigma_{t:t+H} = \frac{1}{K-1} \sum_{k \in \mathcal{K}} |a_{t:t+H}^{(k)} - \mu_{t:t+H}|$ 
     $q(a_{t:t+H}) \leftarrow \text{Normal}(\mu_{t:t+H}, \sigma_{t:t+H}^2 \mathbb{I})$ 
return first action mean  $\mu_t$ .

```

Fig. 47. Latent planning with CEM [20]

#### 4.4.3. Data collection

Since the agent may not initially visit all parts of the environment, new experience needs to be iteratively collect and then the dynamics model gets refined. It is done so by planning with the partially trained model. Starting from a small amount of 6 seed episodes collected under random actions, the model is trained and one additional episode is added to the data set every 5000 update steps.

#### 4.4.4. Preprocessing

Each image gets preprocessed by reducing the bit depth to 5 bits as in [31]. It is then resized to 64 x 64 pixels for all environments. Actions are one-hot encoded and each action gets repeated 4 times as common in reinforcement learning [37] to reduce the planning horizon and provide a clearer learning signal to the model.

#### 4.4.5. Implementation details

This time official code from repository [19] was adjusted and used for experiments. The code architecture follows principles from other PlaNet author's paper [18]. The RSSM uses a GRU [7] with 200 units as deterministic path in the dynamics model and implements all other functions as two fully connected layers of size 200 with ReLU activations [40]. Distributions in latent space are 30-dimensional diagonal Gaussians with predicted mean and standard deviation. The observation model and the approximate state posterior are implemented with the Variational Autoencoder, the same as in World Models (see fig.44). The reward model is implemented with a feed-forward neural network with two hidden layers of size 100. The model is trained jointly using the Adam optimizer [30] with a learning rate of  $10^{-3}$  and  $\epsilon = 10^{-4}$ , and gradient clipping norm of 1000 on batches of 50 sequence chunks of length 50. The KL divergence terms are also scaled relatively to the reconstruction terms and the model is granted free nats by clipping the divergence loss below this value. Both parameters are tuned in the experiments chapter. The latent overshooting from the paper [20] is used with overshooting KL divergence terms additionally scaled by a factor of 1/50. *[QUESTION: Should we rewrite here full loss? It's quite enormous and complicated. Reader can check it in the PlaNet's paper.]*

Planner uses planning horizon of 12, which means that it evaluates 12 actions in the future. Starting from zero mean and unit variance, 1000 candidate action sequences are sampled and evaluated under the learned model. Then the belief gets re-fitted to the top 100 action sequences with the highest fitness scores. After 10 iterations, the planner returns the mean of the belief for the current time step  $\mu_t$  which is then used by the policy to choose discrete action. When collecting episodes for the training data set the epsilon greedy policy is used with  $\epsilon = 0, 3$ . During test phase the greedy policy is used, which chooses the action that maximises the returned belief.

Hiper-parameters presented are used as defaults in experiments described in the next chapter.

## 5. EXPERIMENTS

In this sections the two architectures described in the previous chapter are subject to different experiments that aim at making them to work. The architectures are evaluated using two metrics:

- The more accurate the model at the cut-off point, which is environment dependent, the better model learning algorithm.
- The higher final score of the planning agent using this learned model, the better.

The first metric is evaluated using observations reconstructions at each timestep which are compared to ground truth recordings from the dataset. The second metric is simply final score from the environment. Before experiments descriptions benchmarks and used hardware get reviewed.

### 5.1. Benchmarks

#### 5.1.1. Arcade Learning Environment

*[TODO: Include paragraph about partial observability of ALE.]*

The Arcade Learning Environment (ALE) has became a platform for evaluating artificial intelligence agents. Originally proposed by Bellemare et. al. [2], the ALE makes available dozens of Atari 2600 games for an agent training and evaluation. The agent is expected to do well in as many games as possible without game-specific information, generally perceiving the world through a video stream. Atari 2600 games are excellent environments for evaluating AI agents for three main reasons: they are varied enough to provide multiple different tasks, requiring general competence, they are interesting and challenging for humans and they are free of experimenter's bias, having been developed by an independent party.

In the context of the ALE, a discrete action is a number in range from 0 to 17 inclusive which encodes the composition of a joystick direction and an optional button press. The agent observes a reward signal, which is typically the change in the player's score (the difference in score between the previous time step and the current time step), and an observation  $o_t \in \mathcal{O}$  of the environment. This observation can take form of a single  $210 \times 160$  image and/or the current 1024-bit RAM state. Because a single image typically does not satisfy the Markov property the ALE is formalised as POMDP. Observations and the environment state are distinguished, with the RAM data being the real state of the emulator. A frame (as a unit of time) corresponds to 1/60th of a second, the time interval between two consecutive images rendered to the television screen. The ALE is deterministic, which means that given a particular emulator state  $s$  and a action  $a$  there is a unique next state  $s'$ , that is,  $P_{ss'}^a = p(s'|s, a) = 1$ .

Agents interact with the ALE in an episodic fashion. An episode begins by resetting the environment to its initial configuration,  $s_0$ , and ends at a given endpoint depending on a game. The primary measure of an agent's performance is the score achieved during an episode, namely the undiscounted sum of rewards for that episode. While this performance measure is quite natural, it is important to realize that score is not necessarily an indicator of AI progress. In some

games, agents can exploit the game's mechanics to maximize sum of rewards, but not complete the game's goal in human's understanding [9].

Preprocessing include frame skipping [39] which restricts the agent's decision points by repeating a selected action for 4 consecutive frames. Frame skipping results in a simpler reinforcement learning problem and speeds up execution. *[TODO: Describe other preprocessing techniques used here.]*

This work uses ALE through OpenAI Gym API [4], specifically two games are used as benchmarks: Boxing and Freeway.

Boxing is a video game based on the sport of boxing. Boxing shows a top-down view of two boxers, one white and one black. When close enough, a boxer can hit his opponent with a punch. This causes his opponent to reel back slightly and the boxer scores a point, a reward of 1. In the other situation, when the boxer gets hit, he gets a negative reward of -1. There are no knockdowns or rounds. A match is completed either when one player lands 100 punches (a 'knockout') or two minutes have elapsed. In the case of a decision, the player with the most landed punches is the winner. Ties are possible. While the gameplay is simple, there are subtleties, such as getting an opponent on the 'ropes' and 'juggling' him back and forth between alternate punches.



Fig. 51. Example of Boxing level

In Freeway an agent controls a chicken who can be made to run across a ten lane highway filled with traffic in an effort to "get to the other side." Every time a chicken gets across a reward of 1 is earned by the agent. If hit by a car, then a chicken is forced back slightly. The goal is to score as much points as possible in the two minutes. The chicken is only allowed to move up or down. The major challenge in this environment are sparse rewards. The agent scores only when successfully crosses the highway, which is not a trivial task.



Fig. 52. Example of Freeway level

*[TODO: Add more games if needed.]*

### 5.1.2. Sokoban

*[TODO: Include paragraph about fully observability of Sokoban state.]*

Sokoban is a classic planning problem. It is a challenging one-player puzzle game in which the goal is to navigate a grid world maze and push boxes onto target tiles. A Sokoban puzzle is considered solved when all boxes are positioned on top of target locations. The player can move in all 4 cardinal directions and only push boxes into an empty space (as opposed to pulling). For this reason many moves are irreversible and mistakes can render the puzzle unsolvable. A human player is thus forced to plan moves ahead of time. Artificial agents should similarly benefit from a learned model and simulation.

Despite its simple rule set, Sokoban is an incredibly complex game for which no general solver exists. It can be shown that Sokoban is NP-Hard and PSPACE-complete [10]. Sokoban has an enormous state space that makes it inassailable to exhaustive search methods. An efficient automated solver for Sokoban must have strong heuristics, just as humans utilize their strong intuition, so that it is not overwhelmed by the number of possible game states.

The implementation of Sokoban [46] used for those experiments procedurally generates a new level each episode. This means an agent cannot memorize specific puzzles. Together with the planning aspect, this makes for a very challenging environment. While the underlying game

logic operates in a  $10 \times 10$  grid world, with 7 possible elements in each grid [53], agents were trained directly on RGB sprite graphics. Fig. 54 shows an example of Sokoban level with 4 boxes.

*[TODO: Go into deeper details about e.g. how rewards are obtained etc.]*

| Type       | State      | Graphic |
|------------|------------|---------|
| Wall       | Static     | ■       |
| Floor      | Empty      | ■       |
| Box Target | Empty      | ■       |
| Box        | Off Target | ■       |
| Box        | On Target  | ■       |
| Player     | Off Target | ■       |
| Player     | On Target  | ■       |

Fig. 53. Table with Sokoban possible elements in each grid, further referred to as blocks.

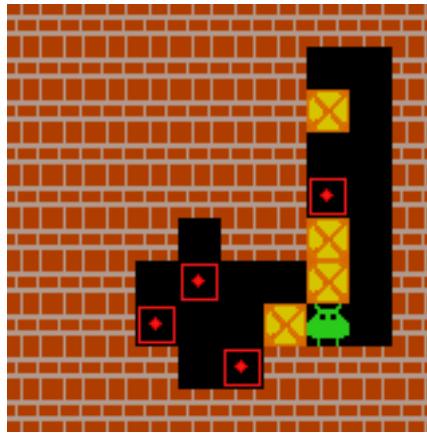


Fig. 54. Example of Sokoban level

## 5.2. Hardware

*[TODO: Add HW specification.]*

## 5.3. World Models for Sokoban

This section focuses on a problem of model learning in Sokoban that is used to solve the environment. The goal was to train a model to generate sharp and accurate future predictions of observations and obtain high score using it. In theory, how probable are sequences of ground truth observations is measured using log probability. In practice, though, it is far more useful to compare those generated sequences of future observations with ground truth sequences with an eye of the researcher. What the researcher looks for are sharp generated images which accurately resemble frames from the game. Moreover, the sequence needs to simulate subsequent actions properly, otherwise it is told that the sequence is noisy or does not model actions issued well.

To generate a future prediction 60 consecutive frames and actions on average are fed into the model to initialize its hidden state first. Then, the model is ready to generate a future prediction with its Memory module by feeding it with subsequent actions and its own predictions as a contemporary state at following time steps. Reader should notice that what gets generated are not frames, but latent state's vectors. Those latent states can then be decoded into images for inspection.

#### 5.3.1. *Train the unchanged World Models implementation in the Sokoban environment*

In this experiment, the original World Models was trained in the Sokoban environment. No modification to the original method was made, beyond addition of the deterministic variant of Memory module described in the previous chapter.

The Vision module successfully learned to encode high dimensional observations into low dimensional latent states. Fig. 55 shows original observations (first and third columns) side by side with reconstructed observations from their encodings (second and fourth columns). These are zero step predictions, no future is predicted only encoding to latent space and decoding to image space again is done.

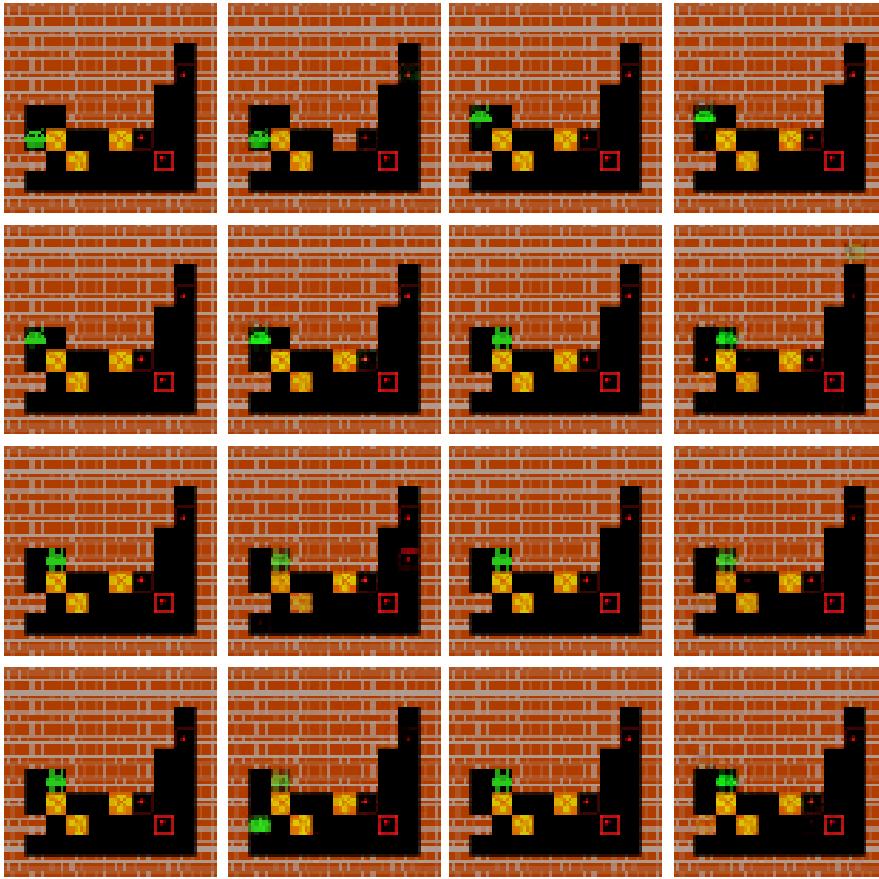


Fig. 55. Qualitative result of the Vision module training in Sokoban. First and third columns include original observations. Second and fourth columns include reconstructions. Each reconstruction was obtained by first encoding the original observation and then decoding it, using VAE encoder and decoder respectively.

The stochastic and deterministic Memory modules were not able to learn Sokoban's dynamics. Fig. 56 shows that the stochastic model very often can not determine the agents position. The agent disappears and blocks change into other blocks. The eighth row shows that pushing mechanics are not modeled, the agent passes through boxes. The deterministic model does not do better. The Controller module failed to learn how to solve any level, it behaved comparably to a random play. We suspect that VAE is unable to generate usable abstract Sokoban representation and the shallow Memory and Controller modules can not grasp complex dynamics of Sokoban using this poor representation. This idea is further developed in the next experiment.

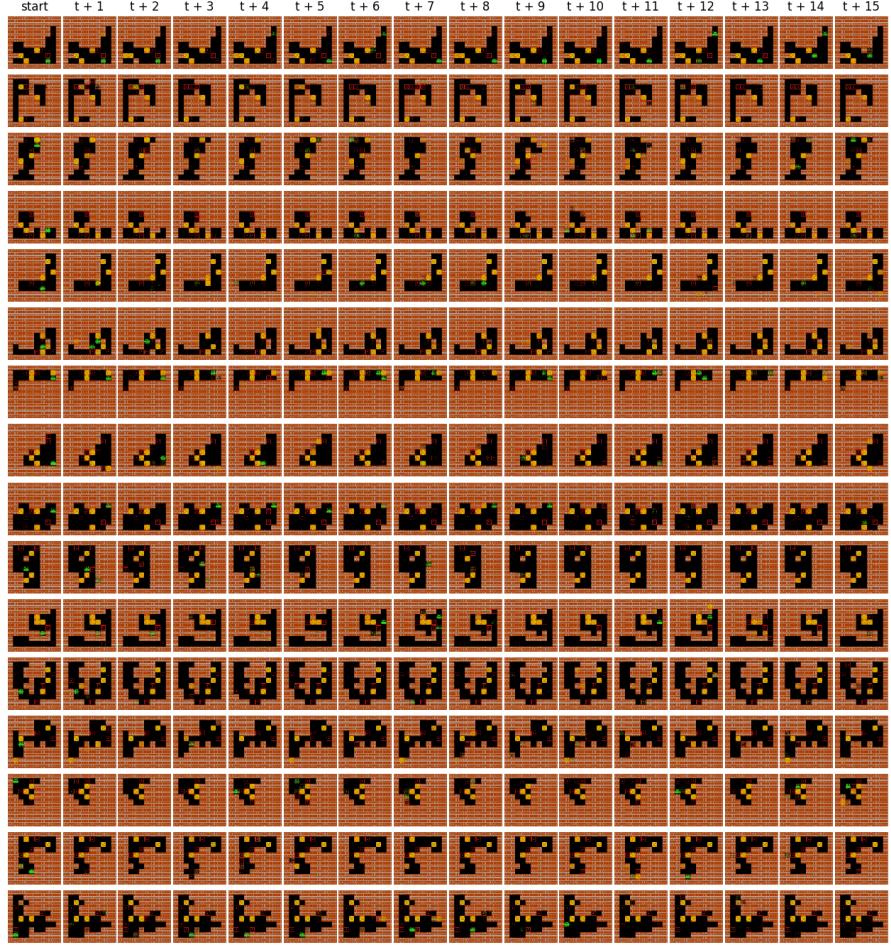


Fig. 56. Qualitative result of the Memory module training in Sokoban. Each row depicts the Memory module rollout in one episode. The first column include original observations from the evaluation dataset from which the rollouts start. The RNN’s hidden state was initialized on preceding transitions in each episode. Each subsequent reconstruction was obtained by first predicting the next latent state by the Memory module and then decoding it using the VAE decoder.

### 5.3.2. Train World Models in Sokoban environment on 10x10 grid world states

The latent state vector size is set to 64. This means that, in theory, this vector can accommodate full information about an observation. As noted before, Sokoban underlying game logic operates in a  $10 \times 10$  grid world, where far edges of a level are always walls. This means that the level is described by 64 blocks organized in an  $8 \times 8$  grid. In this experiment, this domain knowledge is exploited and the agent uses those 64 blocks as an input vector to the Memory module, bypassing the Vision module. It is worth noting, that the Vision module should learn this representation as it is the optimal encoding when the objective is to compress a pixel image into

a 64-dimensional vector and then reconstruct the original observation from it. However, despite use of the optimal encoding, the results have not been improved.

The proposed input format is optimal encoding if one wants to compress a pixel image and then reconstruct it. However, it is really poor representation of a current state of the environment if one wants to use linear combination of those features (blocks in each position) to infer optimal next action and this is exactly what the Controller module is trying to do. Modeling a value function could have more sense e.g. the value function could learn that a box on a target position yields higher value, but even it would have a hard time modeling more complex relations between entities in the environment. More useful for the Controller would be e.g. representation that includes information about distance between the box and each target position. Nevertheless, this could be not enough too. The box on the target position would get discounted for not being on some other target positions. Hence, there is need for feature saying “the box X placed on the target position Y”. In the end, the linear combination of the proposed latent features can not model useful policy.

On the other hand, this representation includes, not well represented, but perfect information about an environment state. The Memory module creates its own environment representation encoded in its hidden state and then uses this representation to predict the next latent state. This Memory’s hidden state is also utilized by the Controller. Still, it does not seem to encode useful enough information for the two, Memory and Controller, to do well on their tasks. One way to improve the hidden state representation is explored in the next experiment.

### 5.3.3. Train World Model in Sokoban with auxiliary tasks

Auxiliary tasks [25] have proved to help create more informative representation of an environment. In this experiment, reward and value prediction tasks are added to the Memory module. In short, two additional linear models are added on top of the RNN to predict the next reward in the environment and model a value function *[TODO: Add information how you train values i.e. Monte-Carlo prediction]*. In theory, it should help form a more informative hidden state of the Memory module. Consequently, it should help learn Sokoban’s dynamics, but also generate representation on a higher level of abstraction that could prove useful for the Controller. Moreover, a reward prediction will be needed in further work on planning with learned model.

For all that, the Memory module has not been able to learn to predict the rewards and values. Also, there was no improvement in Memory’s and Controller’s performance. It is suspected, that the main cause of this failure are sparse rewards in the training dataset. A random agent used to generate the dataset does not receive many positive rewards. Effectively, most of the episodes do not have any positive reward. Hence, the Memory module soon overfit on more or less constant reward and value. This yields insight that the data generation procedure does not cover state-space well. Iterative approach to gathering data, from a better and better agent, could solve this problem.

It is not without significance that Sokoban has enormous state-space. Because each episode, or level, is randomly generated it is much different from the others - it is nearly impos-

sible for an agent to see a similar state in a different episode. Hence, Sokoban requires strong generalization capability from the Memory module. Simple RNN can lack capacity to create good representation and in turn achieve good prediction performance. A more flexible Memory module with larger capacity could manage this complexity and need for generalization. The two insights are explored in the PlaNet experiments, which have larger model and uses iterative training procedure.

#### 5.4. ***World Models for Atari***

Two more experiments below put World Models into the test one more time. Firstly, the original World Models is trained for the Boxing environment, which has dense rewards and data collection using a random agent cover most of the state-space. Then, World Models is coupled with AlphaZero planner and both are trained jointly.

##### 5.4.1. *Train unchanged World Models in the Boxing environment*

In this experiment, the original World Models, with stochastic Memory module, was trained in the Boxing environment. The Vision module successfully learned to encode high dimensional observations into low dimensional latent states. Fig. 57 shows original observations (first and third columns) side by side with reconstructed observations from their encodings (second and fourth columns). These are zero step predictions, no future is predicted only encoding to latent space and decoding to image space again is done.

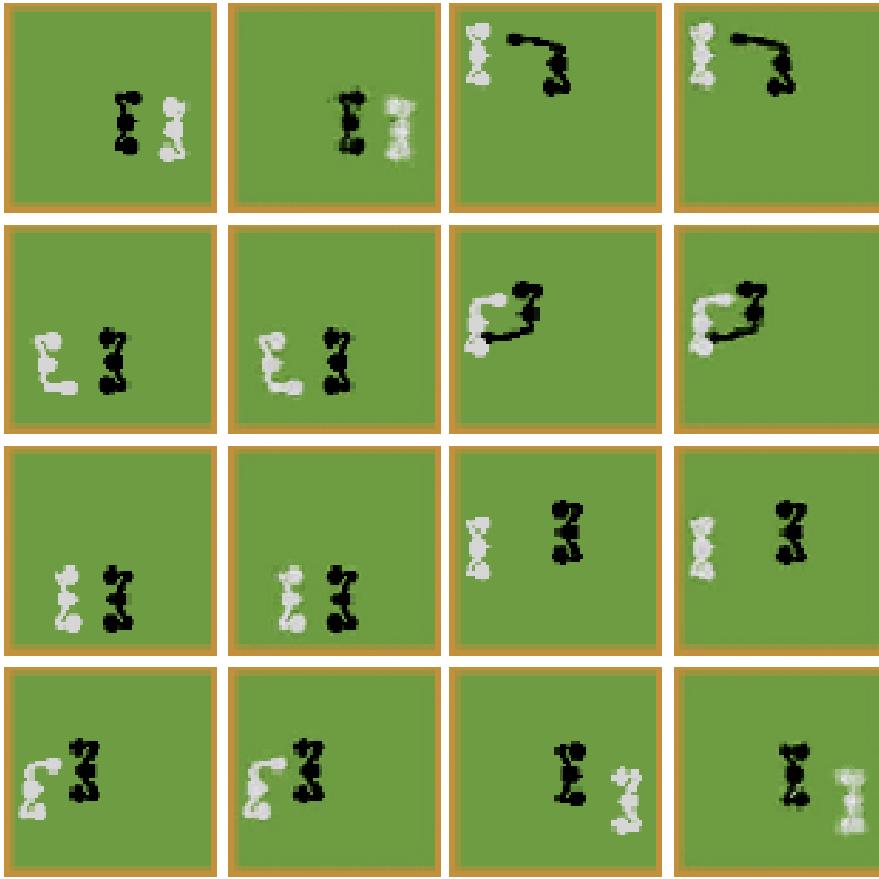


Fig. 57. Qualitative result of the Vision module training in Boxing. First and third columns include original observations. Second and fourth columns include reconstructions. Each reconstruction was obtained by first encoding the original observation and then decoding it, using VAE encoder and decoder respectively.

The stochastic Memory modules was able to learn Boxing’s dynamics. Fig. 58 shows that the stochastic model generates very sharp and accurate predictions that model agents movement and punches really well. The agents does not disappear like in Sokoban and actions are smooth. The Controller module successfully learned how to solve the game scoring above 18 points on average across 5 runs. We suspect that World Models with latent state of size 16 was forced to encode two characters positions and hands states which are useful high-level features when deciding on the next action. It is worth pointing out here that similar experiment with such a small latent space did not yield improvement in Sokoban.

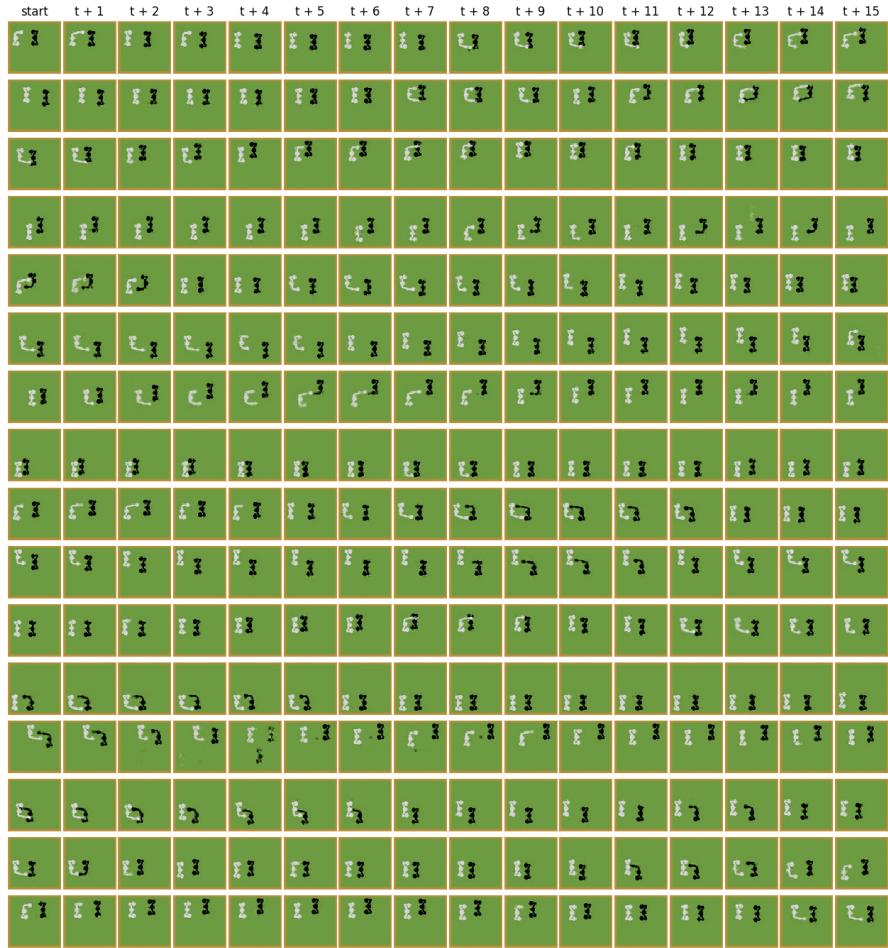


Fig. 58. Qualitative result of the stochastic Memory module training in Boxing. Each row depicts the Memory module rollout in one episode. The first column include original observations from the evaluation dataset from which the rollouts start. The RNN's hidden state was initialized on preceding transitions in each episode. Each subsequent reconstruction was obtained by first predicting the next latent state by the Memory module and then decoding it using the VAE decoder.

#### 5.4.2. Train World Models with AlphaZero planner in the Boxing environment

Despite many attempts and hyper-parameter tuning, the deterministic Memory module was not able to model the Boxing dynamics as good as the stochastic one. This only proves that stochastic nodes are key for the accurate modeling. Fig. 59 shows future predictions which, as can be seen, are imperfect and noisy. The AlphaZero planner training was unstable and it did not train to properly plan using this model. Therefore, it was not able to play the game. Because the AlphaZero planner in its current form can only work with deterministic world models, decision was to abandon this solution and move to the architecture which adjusts PlaNet, which shown that, in

deed, it is possible to plan in continuous control tasks using a stochastic model with iterative data collection procedure.

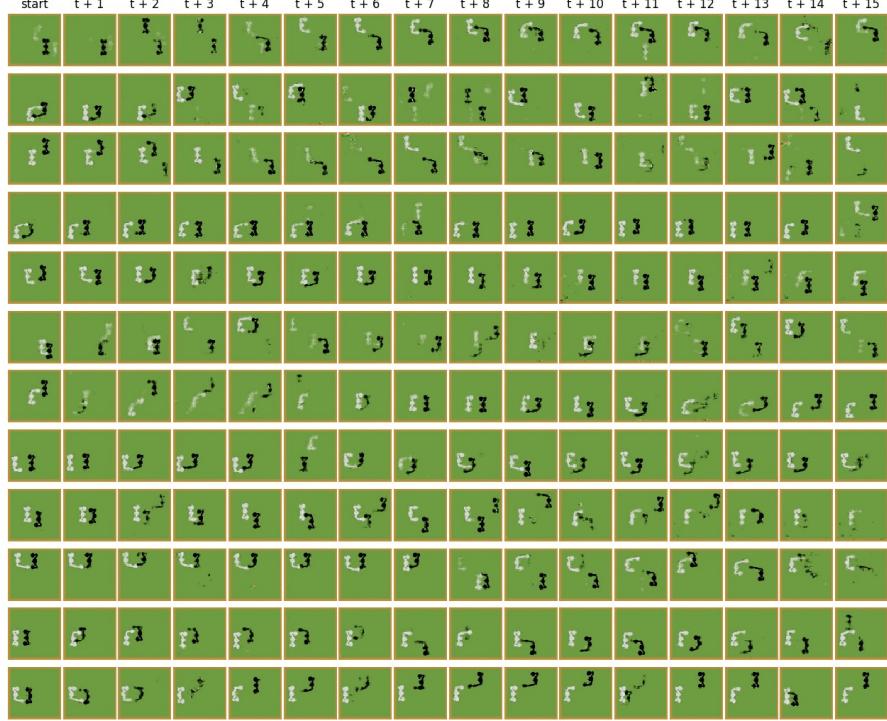


Fig. 59. Qualitative result of the deterministic Memory module training in Boxing. Each row depicts the Memory module rollout in one episode. The first column include original observations from the evaluation dataset from which the rollouts start. The RNN's hidden state was initialized on preceding transitions in each episode. Each subsequent reconstruction was obtained by first predicting the next latent state by the Memory module and then decoding it using the VAE decoder.

## 5.5. PlaNet for Sokoban

### 5.5.1. Train unchanged PlaNet in the Sokoban environment

PlaNet did not capture Sokoban dynamics too. In the figure below (Fig. 510) future predictions are blurred, multiple agents appear and other artifacts, like changing blocks, are present. Similarly like in World Models case decision was to move to Atari games as easier environments to start with.

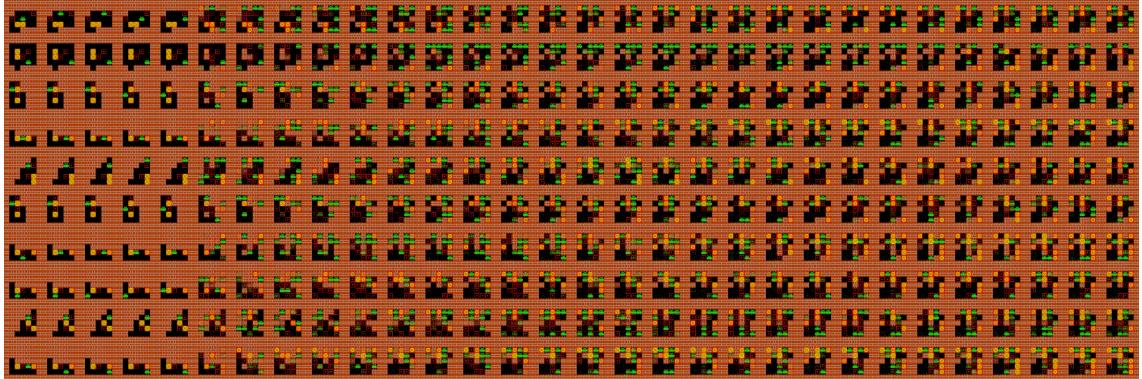


Fig. 510. Qualitative result of the model training in Sokoban. Each row depicts the model rollout in one episode. The first five columns include original consecutive observations from the evaluation dataset from which the rollouts start. The model hidden state was initialized on these transitions. Each subsequent reconstruction was obtained by first predicting the next latent state by the model and then decoding it using the decoder.

## 5.6. PlaNet for Atari

In this section experiments that lead to first successful case are described. The next section will focus on tuning this method to yield high scores, comparable to model-free methods, with the smallest amount of data possible.

### 5.6.1. Train unchanged PlaNet in the Boxing environment

It did not start to work out of the box of course. Fig. 511 shows that future predictions turn into a blurry blob, where it is not possible to distinguish one player from another.

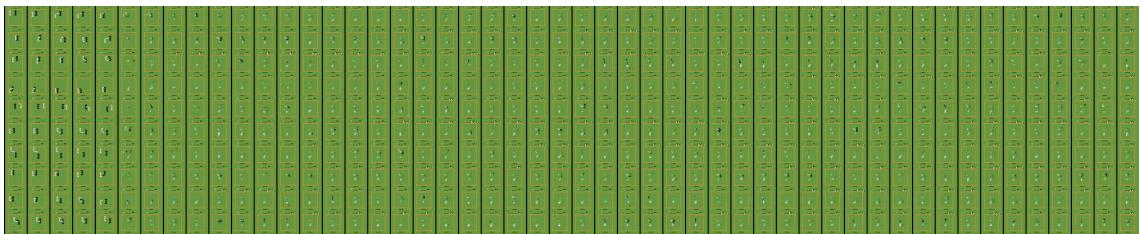


Fig. 511. Qualitative result of the model training in Boxing. Each row depicts the model rollout in one episode. The first five columns include original consecutive observations from the evaluation dataset from which the rollouts start. The model hidden state was initialized on these transitions. Each subsequent reconstruction was obtained by first predicting the next latent state by the model and then decoding it using the decoder.

By default a decoder variance of 1 is used, which means the model explains a lot of variation in the image as random noise. While this leads to more robust representations, it also leads to more blurry images. If the changes in consecutive frames are minor, then the posterior collapses because the model explains everything as observations noise. There are two possible solutions to this issue: one is to increase an action repeat and the other is to try to reduce the decoder variance. These are examined next.

### 5.6.2. Train PlaNet in the Boxing environment with the increased action repeat

The action repeat will result in a bigger difference between consecutive frames and thus more signal for the model to learn from, that cannot be easily modeled as noise. In practice though, it did not help and even made the agent play worse than a random agent. The random agent is taking moves at random.

### 5.6.3. Train PlaNet in the Boxing environment with the lowered decoder variance

The predictions are more blurry with a higher variance because the decoder generate more observations that differ slightly from the same latent code. This leads to the posterior explaining more similar observations with the same code. If consecutive frames are very similar, then the posterior collapses and explain them with one code. By lowering the variance it becomes more sensitive to small changes in observations. *[TODO: Proofread this explanation once more and/or ask Danijar if it is right.]*

*[TODO: Move implementation details to project's plan and refer to them here.]* Lowering the decoder variance is equivalent to lowering a KL divergence scale in the PlaNet loss. It can be seen by writing the ELBO for a Gaussian decoder in the standard form  $E_q(z)[\ln p(x|z)] - KL[q(z)||p(z)]$ . The log-likelihood terms is  $\ln p(x|z) = -0.5(x - f(z))/\sigma^2 - \ln Z$ . Multiplying the ELBO by  $\sigma^2$  removes it from the log-probability term and puts it in front of the KL term as in beta-VAE [23]. The objectives have different values because of the Gaussian normalizer  $Z$  but they share the same gradient since the normalizer is a constant. Other reason that lowering the divergence scale can help with collapsing posterior is that it allows the model to absorb more information from its observations by loosening the information bottleneck.

On the other hand it is recommended to keep the divergence scale as high as possible while still allowing for good performance. For example, when the divergence scale is set to zero it could learn to become a deterministic autoencoder which reconstruct observations well but is less likely to generalize to state in latent space that the decoder hasn't seen during training.

Random search resulted in the best divergence scale being around 0.03. It was tuned jointly with a free nats parameter which is described in the next section. *[QUESTION: I should add diagram with random search results, but how to do this if those are evaluated with a researcher eye?] [TODO: You should better describe this random search experiment. What parameters where tuned, which turned out to be the most important, for how long and how much runs you were running etc.]*

### 5.6.4. Train PlaNet in the Boxing environment with increased free nats

*[TODO: Move theoretical details to theoretical background and refer to them here.]* Free nats technique is often used for static Variational Autoencoders. The model is allowed to use given amount of nats without KL penalty in the variational objective. It helps the model focus on smaller details which do not contribute much to improving the reconstruction loss. Intuitively to this threshold of KL divergence (between a prior and a posterior) reconstruction loss is favoured.

In case of Boxing, it helped to model boxers moves and actions more accurately. The best free nats turned out to be 12. Fig. 512 shows final result.

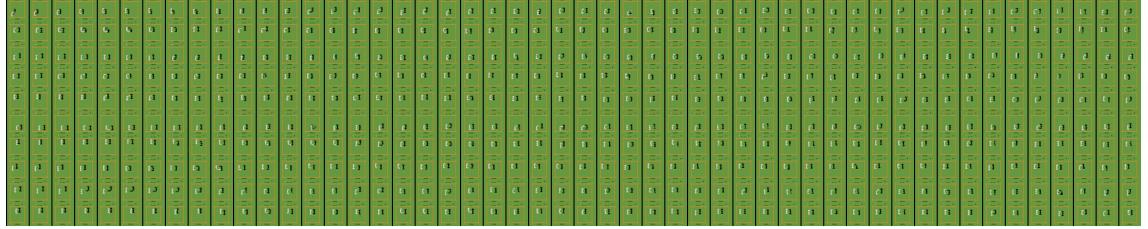


Fig. 512. Qualitative result of the model training in Boxing. Each row depicts the model rollout in one episode. The first five columns include original consecutive observations from the evaluation dataset from which the rollouts start. The model hidden state was initialized on these transitions. Each subsequent reconstruction was obtained by first predicting the next latent state by the model and then decoding it using the decoder.

It achieved final score around 45 after 1 million steps in the environment. The learning curve is shown in fig. 513.

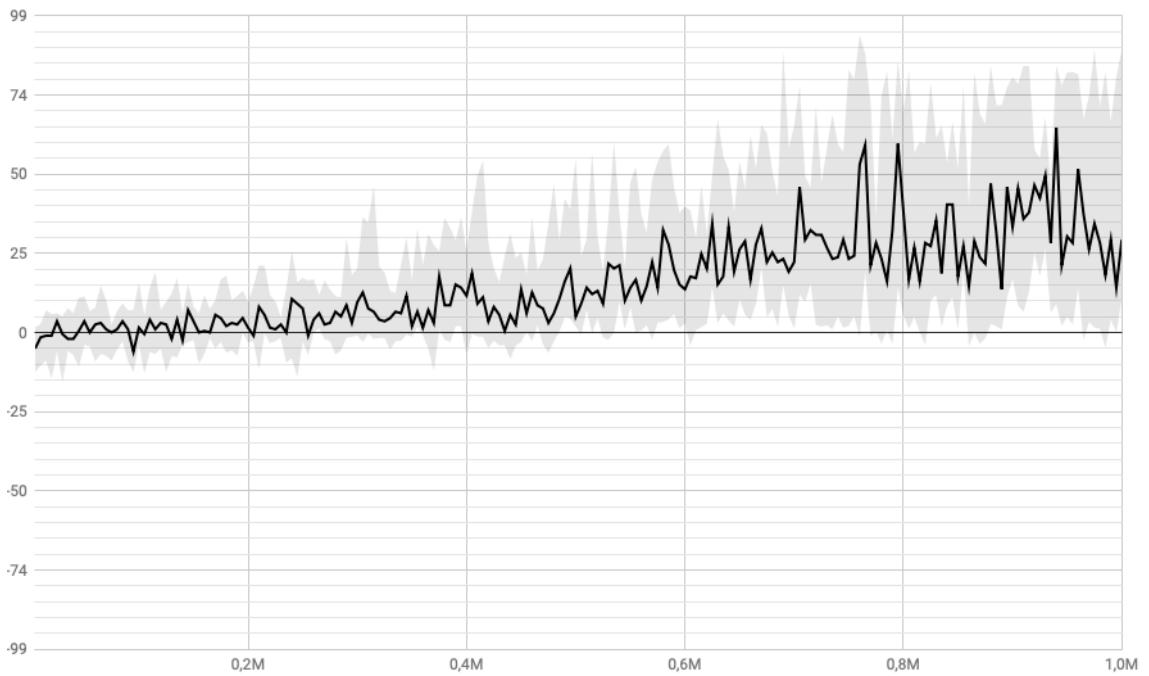


Fig. 513. PlaNet learning curve after tuning and with overshooting. The line presents median and the shaded area are percentiles 5 to 95 over 5 training runs. On the Y axis is the environment score and on the X axis is number of steps in the real environment for the data collection.

Table. 51 summarize hyper-parameters tuning experiment. The specialist evaluated noisiness of the model rollouts, the same as in fig. 512, in scale from 0 to 2, where 0 is a little or no noise in the future reconstructions and 2 is a lot of noise in the future reconstructions.

| Free nats | Divergence scale | Noisiness |
|-----------|------------------|-----------|
| 2         | 1E-04            | 2         |
| 16        | 1E-04            | 2         |
| 13        | 1E-04            | 2         |
| 11        | 2E-04            | 2         |
| 3         | 2E-04            | 2         |
| 1         | 3E-04            | 2         |
| 8         | 3E-04            | 2         |
| 4         | 4E-04            | 2         |
| 17        | 1E-03            | 2         |
| 7         | 2E-03            | 2         |
| 10        | 2E-03            | 2         |
| 15        | 2E-03            | 2         |
| 6         | 2E-03            | 2         |
| 9         | 3E-03            | 2         |
| 13        | 4E-03            | 2         |
| 15        | 4E-03            | 2         |
| 15        | 5E-03            | 2         |
| 8         | 5E-03            | 2         |
| 5         | 6E-03            | 2         |
| 10        | 1E-02            | 1         |
| 4         | 2E-02            | 1         |
| 8         | 2E-02            | 0         |
| 17        | 3E-02            | 0         |
| 12        | 1E-01            | 0         |

Table 51. PlaNet for Boxing hyper-parameters tuning results.

The lower divergence scale the nosier predictions are. The higher free nats the better, more stable, movement predictions are. Best params turned out to be: the divergence scale around 3E-02 and the free nats around 12.

#### 5.6.5. Train tuned PlaNet in the Freeway environment

The same random search procedure was applied to the Freeway environment described earlier. Fig. 514 shows final result.

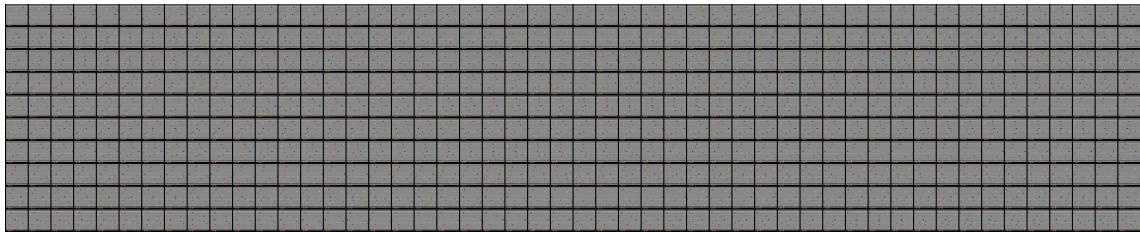


Fig. 514. Qualitative result of the model training in Freeway. Each row depicts the model rollout in one episode. The first five columns include original consecutive observations from the evaluation dataset from which the rollouts start. The model hidden state was initialized on these transitions. Each subsequent reconstruction was obtained by first predicting the next latent state by the model and then decoding it using the decoder.

Table. 52 summarize hyper-parameters tuning experiment in the same way like for Boxing.

| Free nats | Divergence scale | Noisiness |
|-----------|------------------|-----------|
| 9         | 8E-02            | 2         |
| 6         | 2E-02            | 2         |
| 5         | 3E-03            | 2         |
| 4         | 6E-04            | 2         |
| 4         | 4E-04            | 2         |
| 4         | 2E-04            | 2         |
| 7         | 3E-02            | 1         |
| 6         | 5E-04            | 1         |
| 4         | 1E-04            | 1         |
| 2         | 2E-04            | 1         |
| 5         | 8E-03            | 0         |
| 2         | 2E-03            | 0         |

Table 52. PlaNet for Freeway hyper-parameters tuning results.

High free nats, above 6, and too low, below 1E-03, or too high, above 9E-03, divergence scale makes future predictions very noisy and blurry. Best parameters chosen were: the divergence scale of 8E-03 and the free nats of 3.

Despite really good future observations prediction the agent failed to solve the task. *[TODO: In benchmark description you should write what you consider as solved task in each environment.]* Possibly planner horizon is to short to cover a plan which ends with positive reward on the other side of the road. This is explored in the next experiment.

#### 5.6.6. Train tuned PlaNet in the Freeway environment with a longer planning horizon

Here two experiments were run: first for planning horizon of 25 and second for planning horizon of 50. Both failed.

#### 5.6.7. Train PlaNet in the Boxing environment without overshooting

The authors in the final version of the PlaNet paper find out, that disabling overshooting can increase performance for their RSSM architecture. It was put into the test and final results are shown in fig. 515. It helped increase performance of RSSM tuned for Boxing.

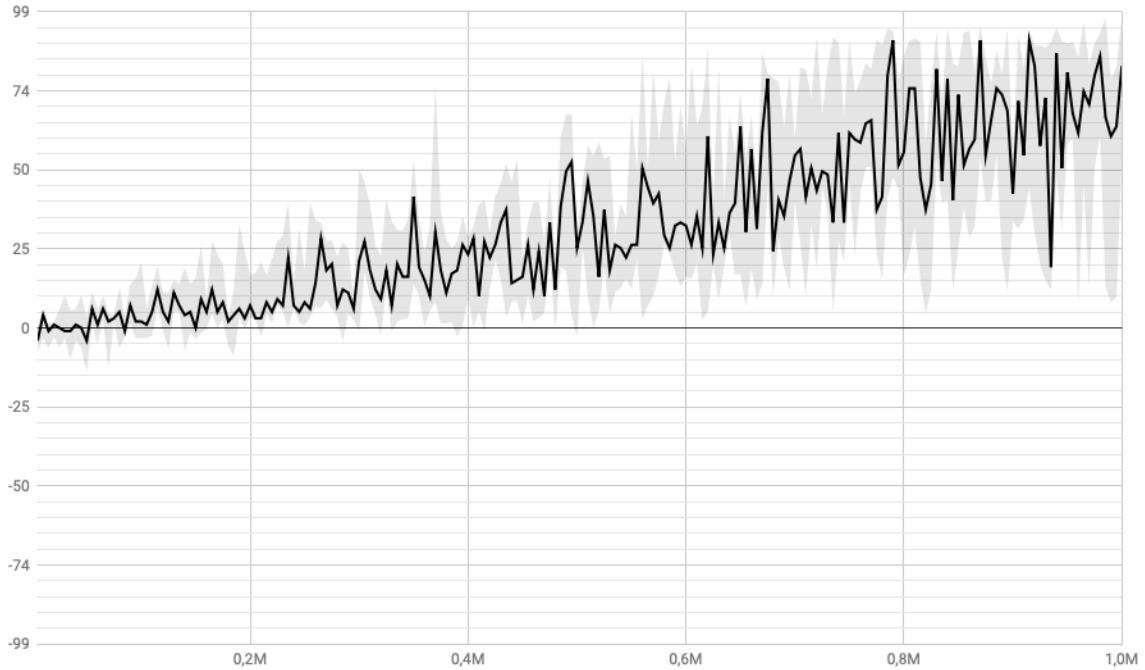


Fig. 515. PlaNet learning curve after tuning and without overshooting. The line presents median and the shaded area are percentiles 5 to 95 over 5 training runs. On the Y axis is the environment score and on the X axis is number of steps in the real environment for the data collection.

This is the best and final solution which gets compared to strong model-based baseline SimPLe [29] and model-free baselines Rainbow [22] and PPO [47]. Comparison is done in low data regime of 100K, 500K and 1M interactions with the real environment. All of the results are taken from the SimPLe paper [29].

| Algo.   | 100K       | 500K        | 1M          |
|---------|------------|-------------|-------------|
| Ours    | 6,2 (10,7) | 35,2 (8,3)  | 78,2 (19,1) |
| SimPLe  | 9,1 (8,8)  | NDA         | NDA         |
| PPO     | -3,9 (6,4) | 3,5 (3,5)   | 19,6 (20,9) |
| Rainbow | 0,9 (1,7)  | 58,2 (16,5) | 80,3 (5,6)  |
| Random  |            | 0,3         |             |

Table 53. Mean scores and standard deviations (in brackets) over five training runs.

Tuned PlaNet architecture without overshooting is better from PPO in every case and from Rainbow in low data regime of 100K interactions with the real environment. It is then comparable to SimPLe and Rainbow in the other cases. It is clear that this architecture is sample-efficient without sacrificing final performance.

## **6. CONCLUSION**

## REFERENCES

- [1] Gabriel Barth-Maron et al. “Distributed Distributional Deterministic Policy Gradients”. In: *arXiv e-prints* (2018). arXiv: 1804.08617.
- [2] Marc G. Bellemare et al. “The Arcade Learning Environment: An Evaluation Platform for General Agents”. In: *arXiv e-prints* (2012). arXiv: 1207.4708.
- [3] Zdravko I. Botev et al. “Chapter 3 - The Cross-Entropy Method for Optimization”. In: *Handbook of Statistics*. Vol. 31. Handbook of Statistics. Elsevier, 2013, pp. 35 –59. DOI: <https://doi.org/10.1016/B978-0-444-53859-8.00003-5>.
- [4] Greg Brockman et al. *OpenAI Gym*. 2016. arXiv: 1606.01540.
- [5] Lars Buesing et al. “Learning and Querying Fast Generative Models for Reinforcement Learning”. In: *arXiv e-prints* (2018). arXiv: 1802.03006.
- [6] Silvia Chiappa et al. “Recurrent Environment Simulators”. In: *arXiv e-prints* (2017). arXiv: 1704.02254.
- [7] Kyunghyun Cho et al. “Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation”. In: *arXiv e-prints* (2014). arXiv: 1406.1078.
- [8] François Chollet et al. *Keras*. <https://keras.io>. 2015.
- [9] Jack Clark and Dario Amodei. *Faulty Reward Functions in the Wild*. <https://openai.com/blog/faulty-reward-functions/>. Accessed: 2019-05-25.
- [10] Dorit Dor and Uri Zwick. “SOKOBAN and other motion planning problems”. In: *Computational Geometry* (1999). ISSN: 0925-7721. DOI: [https://doi.org/10.1016/S0925-7721\(99\)00017-6](https://doi.org/10.1016/S0925-7721(99)00017-6). URL: <http://www.sciencedirect.com/science/article/pii/S0925772199000176>.
- [11] Richard Evans et al. “De novo structure prediction with deep-learning based scoring”. Dec. 2018. URL: <https://deepmind.com/blog/alphafold/>.
- [12] Vladimir Feinberg et al. “Model-Based Value Estimation for Efficient Model-Free Reinforcement Learning”. In: *arXiv e-prints* (2018). arXiv: 1803.00101.
- [13] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. <http://www.deeplearningbook.org>. MIT Press, 2016.
- [14] Alex Graves. “Generating Sequences With Recurrent Neural Networks”. In: *arXiv e-prints* (2013). arXiv: 1308.0850.
- [15] Karol Gregor et al. “Temporal Difference Variational Auto-Encoder”. In: *arXiv e-prints* (2018). arXiv: 1806.03107.
- [16] David Ha and Douglas Eck. “A Neural Representation of Sketch Drawings”. In: *arXiv e-prints* (2017). arXiv: 1704.03477.
- [17] David Ha and Jürgen Schmidhuber. “World Models”. In: *arXiv e-prints* (2018). arXiv: 1803.10122.
- [18] Danijar Hafner, James Davidson, and Vincent Vanhoucke. “TensorFlow Agents: Efficient Batched Reinforcement Learning in TensorFlow”. In: *arXiv e-prints* (2017). arXiv: 1709.02878. URL: <http://arxiv.org/abs/1709.02878>.
- [19] Danijar Hafner et al. *Deep Planning Network*. <https://github.com/google-research/planet>. 2019.
- [20] Danijar Hafner et al. “Learning Latent Dynamics for Planning from Pixels”. In: *arXiv e-prints* (2018). arXiv: 1811.04551v4.
- [21] Nikolaus Hansen. “The CMA Evolution Strategy: A Tutorial”. In: *arXiv e-prints* (2016). arXiv: 1604.00772.
- [22] Matteo Hessel et al. “Rainbow: Combining Improvements in Deep Reinforcement Learning”. In: *arXiv e-prints* (2017). arXiv: 1710.02298.
- [23] Irina Higgins et al. “beta-VAE: Learning Basic Visual Concepts with a Constrained Variational Framework”. In: *ICLR* (2017).

- [24] Sepp Hochreiter and Jürgen Schmidhuber. “Long Short-Term Memory”. In: *Neural Comput.* (1997). ISSN: 0899-7667. DOI: 10.1162/neco.1997.9.8.1735. URL: <http://dx.doi.org/10.1162/neco.1997.9.8.1735>.
- [25] Max Jaderberg et al. “Reinforcement Learning with Unsupervised Auxiliary Tasks”. In: *arXiv e-prints* (2016). arXiv: 1611.05397.
- [26] Se Won Jang, Jesik Min, and Chan Lee. “Reinforcement Car Racing with A3C”. In: (2017). URL: <https://www.scribd.com/document/358019044/Reinforcement-Car-Racing-with-A3C>.
- [27] Piotr Januszewski, Grzegorz Beringer, and Mateusz Jablonski. *HumbleRL - Straightforward reinforcement learning Python framework*. 2019. URL: <https://github.com/piojanu/humblerl>.
- [28] Danilo Jimenez Rezende, Shakir Mohamed, and Daan Wierstra. “Stochastic Backpropagation and Approximate Inference in Deep Generative Models”. In: *arXiv e-prints* (2014). arXiv: 1401.4082.
- [29] Lukasz Kaiser et al. “Model-Based Reinforcement Learning for Atari”. In: *arXiv e-prints* (2019). arXiv: 1903.00374.
- [30] Diederik P. Kingma and Jimmy Ba. “Adam: A Method for Stochastic Optimization”. In: *arXiv e-prints* (2014). arXiv: 1412.6980.
- [31] Diederik P. Kingma and Prafulla Dhariwal. “Glow: Generative Flow with Invertible 1x1 Convolutions”. In: *arXiv e-prints* (2018). arXiv: 1807.03039.
- [32] Levente Kocsis and Csaba Szepesvári. “Bandit Based Monte-Carlo Planning”. In: *Machine Learning: ECML 2006*. Ed. by Johannes Fürnkranz, Tobias Scheffer, and Myra Spiliopoulou. Springer Berlin Heidelberg, 2006.
- [33] Rahul G. Krishnan, Uri Shalit, and David Sontag. “Deep Kalman Filters”. In: *arXiv e-prints* (2015). arXiv: 1511.05121.
- [34] Felix Leibfried, Nate Kushman, and Katja Hofmann. “A Deep Learning Approach for Joint Video Frame and Reward Prediction in Atari Games”. In: *arXiv e-prints* (2016). arXiv: 1611.07078.
- [35] Marlos C. Machado et al. “Revisiting the Arcade Learning Environment: Evaluation Protocols and Open Problems for General Agents”. In: *arXiv e-prints* (2017). arXiv: 1709.06009.
- [36] Volodymyr Mnih et al. “Asynchronous Methods for Deep Reinforcement Learning”. In: *arXiv e-prints* (2016). arXiv: 1602.01783.
- [37] Volodymyr Mnih et al. “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540 (Feb. 2015), pp. 529–533. URL: <http://dx.doi.org/10.1038/nature14236>.
- [38] S Mor-Yosef et al. “Ranking the risk factors for cesarean: Logistic regression analysis of a nationwide study”. In: *Obstetrics and gynecology* 75 (July 1990), pp. 944–7.
- [39] Y. Naddaf and University of Alberta. Department of Computing Science. *Game-independent AI Agents for Playing Atari 2600 Console Games*. University of Alberta, 2010. URL: <https://books.google.pl/books?id=c85vnQAACAAJ>.
- [40] Vinod Nair and Geoffrey E. Hinton. “Rectified Linear Units Improve Restricted Boltzmann Machines”. In: ICML’10 (2010), pp. 807–814. URL: <http://dl.acm.org/citation.cfm?id=3104322.3104425>.
- [41] Junhyuk Oh, Satinder Singh, and Honglak Lee. “Value Prediction Network”. In: *arXiv e-prints* (2017). arXiv: 1707.03497.
- [42] Adam Paszke et al. “Automatic Differentiation in PyTorch”. In: *NIPS Autodiff Workshop*. 2017.
- [43] Sumit Saha. *A Comprehensive Guide to Convolutional Neural Networks*. [Online; accessed 25-June-2019]. 2018. URL: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>.
- [44] Sumit Saha. *Understanding LSTM Networks*. [Online; accessed 25-June-2019]. 2015. URL: <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>.

- [45] Tim Salimans et al. “Evolution Strategies as a Scalable Alternative to Reinforcement Learning”. In: *arXiv e-prints* (2017). arXiv: 1703.03864.
- [46] Max-Philipp B. Schrader. *gym-sokoban*. <https://github.com/mpSchrader/gym-sokoban>. 2018.
- [47] John Schulman et al. “Proximal Policy Optimization Algorithms”. In: *arXiv e-prints* (2017). arXiv: 1707.06347.
- [48] David Silver, Richard S. Sutton, and Martin Müller. “Temporal-difference search in computer Go”. In: *Machine Learning* 87.2 (2012), pp. 183–219. ISSN: 1573-0565. DOI: 10.1007/s10994-012-5280-0. URL: <https://doi.org/10.1007/s10994-012-5280-0>.
- [49] David Silver et al. “A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play”. In: *Science* 362 (2018).
- [50] David Silver et al. “Mastering the game of Go without human knowledge”. In: *Nature* 550 (2017). URL: <http://dx.doi.org/10.1038/nature24270>.
- [51] Ilya Sutskever et al. “On the importance of initialization and momentum in deep learning”. In: (2013). URL: <http://proceedings.mlr.press/v28/sutskever13.html>.
- [52] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction* 2nd. The MIT Press, 2018. ISBN: 9780262039246.
- [53] Erik Talvitie. “Agnostic System Identification for Monte Carlo Planning”. In: AAAI’15 (2015), pp. 2986–2992. URL: <http://dl.acm.org/citation.cfm?id=2888116.2888132>.
- [54] Erik Talvitie. “Model Regularization for Stable Sample Rollouts”. In: UAI’14 (2014), pp. 780–789. URL: <http://dl.acm.org/citation.cfm?id=3020751.3020832>.
- [55] Théophane Weber et al. “Imagination-Augmented Agents for Deep Reinforcement Learning”. In: *arXiv e-prints* (2017). arXiv: 1707.06203.
- [56] Wikipedia contributors. *Ms. Pac-Man — Wikipedia, The Free Encyclopedia*. [Online; accessed 21-June-2019]. 2019. URL: [https://en.wikipedia.org/w/index.php?title=Ms.\\_Pac-Man&oldid=900278576](https://en.wikipedia.org/w/index.php?title=Ms._Pac-Man&oldid=900278576).

## LIST OF FIGURES

|      |   |    |
|------|---|----|
| 21.  | Reinforcement Learning .....  | 11 |
| 22.  | General Policy Iteration [52] .....   | 14 |
| 23.  | Model-based Reinforcement Learning. [52] Each arrow shows a relationship of influence and presumed improvement..... | 16 |
| 24.  | Deep Learning .....   | 19 |
| 25.  | Multilayer perceptron .....   | 20 |
| 26.  | A recurrent neural network [44] .....   | 20 |
| 27.  | A convolutional neural network.....   | 21 |
| 31.  | Fast Generative Models.....   | 24 |
| 32.  | Flow diagram of the World Models agent's model .....  | 27 |
| 33.  | VizDoom .....   | 29 |
| 34.  | PlaNet latent dynamics model designs.....   | 30 |
| 35.  | PlaNet latent dynamics model unrolling schemes .....  | 31 |
| 36.  | DeepMind Control Suite.....   | 32 |
| 41.  | HumbleRL architecture .....   | 37 |
| 42.  | HumbleRL loop function overview .....   | 38 |
| 43.  | World Models probabilistic graphical model .....  | 40 |
| 44.  | World Models VAE neural network architecture [17].....  | 42 |
| 45.  | Monte-Carlo tree search in AlphaZero [50]......   | 44 |
| 46.  | World Models probabilistic graphical model .....  | 46 |
| 47.  | Latent planning with CEM [20].....  | 47 |
| 51.  | Boxing .....  | 50 |
| 52.  | Freeway .....   | 51 |
| 53.  | Sokoban elements .....  | 52 |
| 54.  | Sokoban.....  | 52 |
| 55.  | Qualitative result of the World Models' Vision module training in Sokoban .....                                     | 54 |
| 56.  | Qualitative result of the World Models' Memory module training in Sokoban.....                                      | 55 |
| 57.  | Qualitative result of the World Models' Vision module training in Boxing .....                                      | 58 |
| 58.  | Qualitative result of the World Models' stochastic Memory module training in Boxing                                 | 59 |
| 59.  | Qualitative result of the World Models' deterministic Memory module training in<br>Boxing .....                     | 60 |
| 510. | Qualitative result of the PlaNet model training in Sokoban.....   | 61 |
| 511. | Qualitative result of the original PlaNet model training in Boxing .....  | 61 |
| 512. | Qualitative result of the PlaNet model training with a lower divergence scale in<br>Boxing .....                    | 63 |
| 513. | PlaNet learning curve after tuning and with overshooting .....  | 63 |
| 514. | Qualitative result of the PlaNet model training with a lower divergence scale in<br>Freeway .....                   | 64 |
| 515. | PlaNet learning curve after tuning and without overshooting .....   | 66 |

## **LIST OF TABLES**

|     |  |    |
|-----|--|----|
| 51. | PlaNet for Boxing hyper-parameters tuning results .....  | 64 |
| 52. | PlaNet for Freeway hyper-parameters tuning results ..... | 65 |
| 53. | Results comparison .....                                 | 66 |